

# Deep Learning for beginners

## Part II - ConvNet

**JB Fiche**, CBS-Montpellier & Plateforme MARS-MRI

**Francesco Pedaci**, CBS-Montpellier

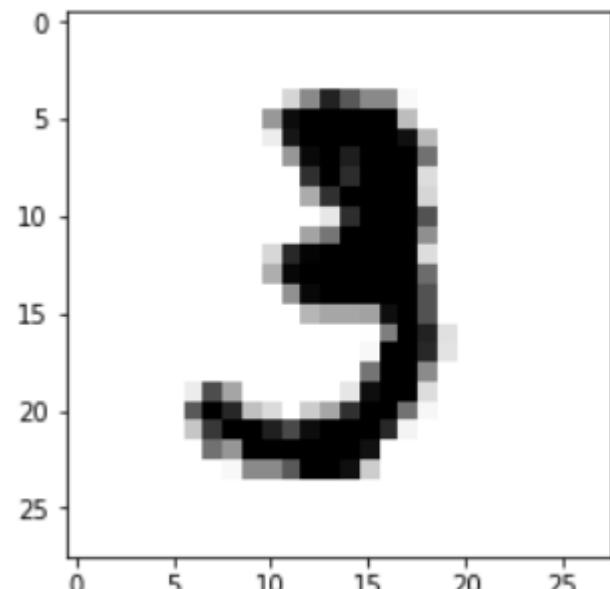
**Volker Bäcker**, CRBM & MRI

**Cédric Hassen-Khodja**, CRBM & MRI

# Outline

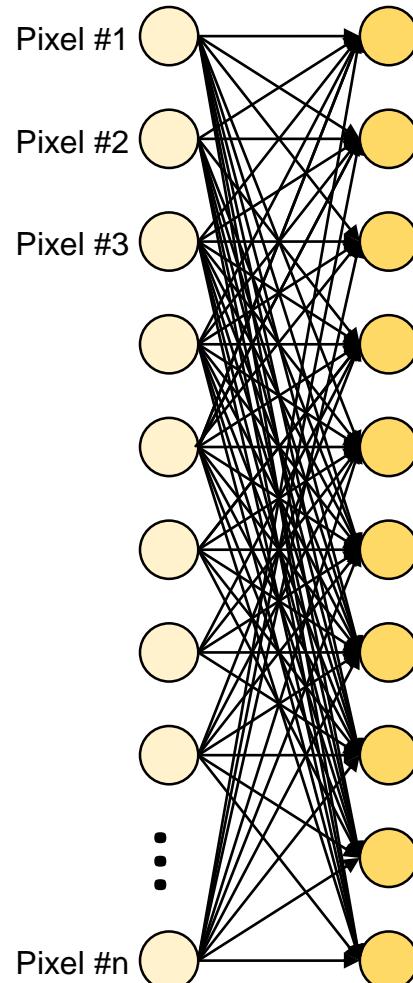
- I. Introduction**
- II. Starting with Deep Learning and Artificial Neuron Network**
- III. Deep Learning for image analysis**
  - i. Introduction to convolutional network**
  - ii. Application : Digit recognition
  - iii. Practical case 1 : image classification
  - iv. Practical case 2 : 2D regression of bead positions
  - v. Practical case 3 : Semantic segmentation of bacteria

# Analyze images with densely connected networks



28x28 pixels image

784 inputs

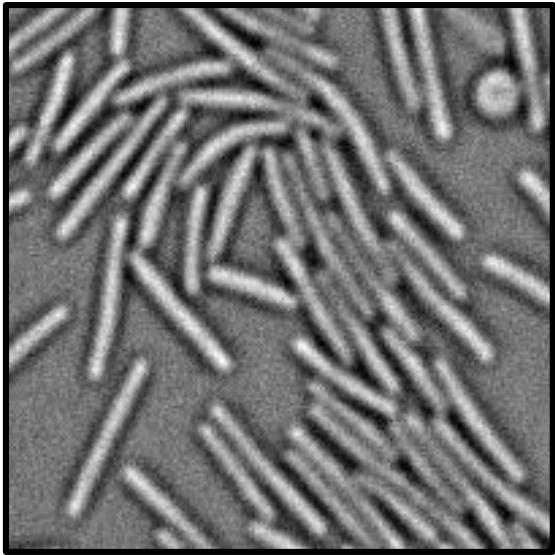


First layer 10  
neurons

**Number of parameters:**

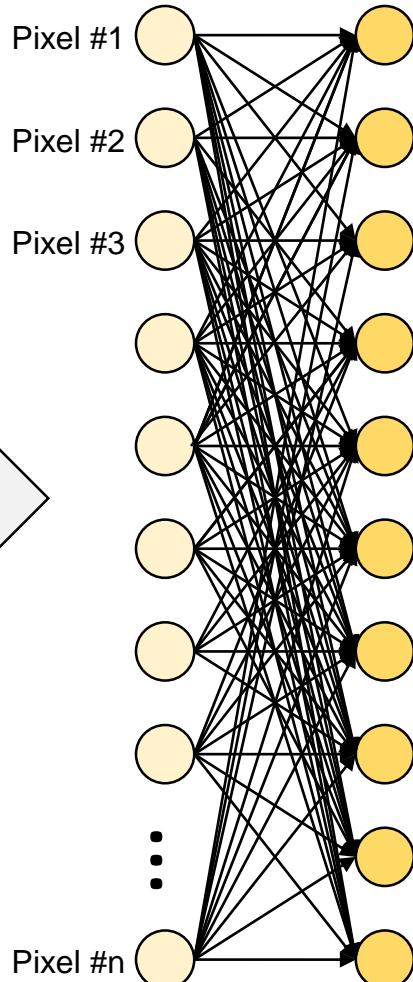
$$784 \times 10 + 10 = 7850$$

# Analyze images with densely connected networks



200x200 pixels image

40,000 inputs



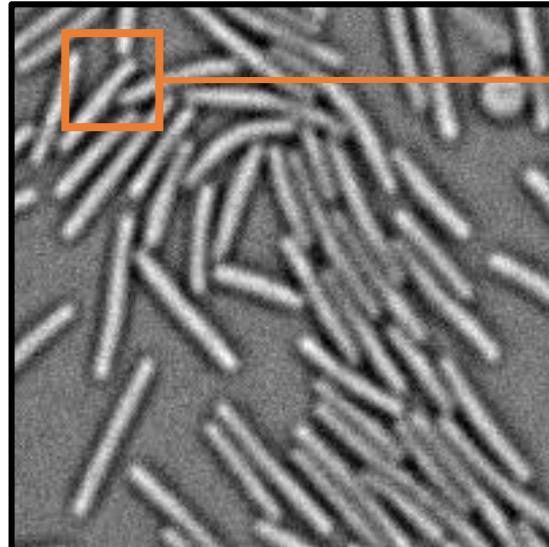
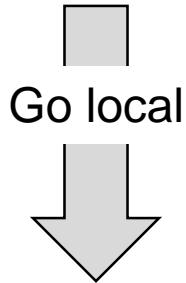
**Number of parameters:**

$$40,000 * 10 + 10 = 400,010 \quad (!!!)$$

# Analyze images with densely connected networks

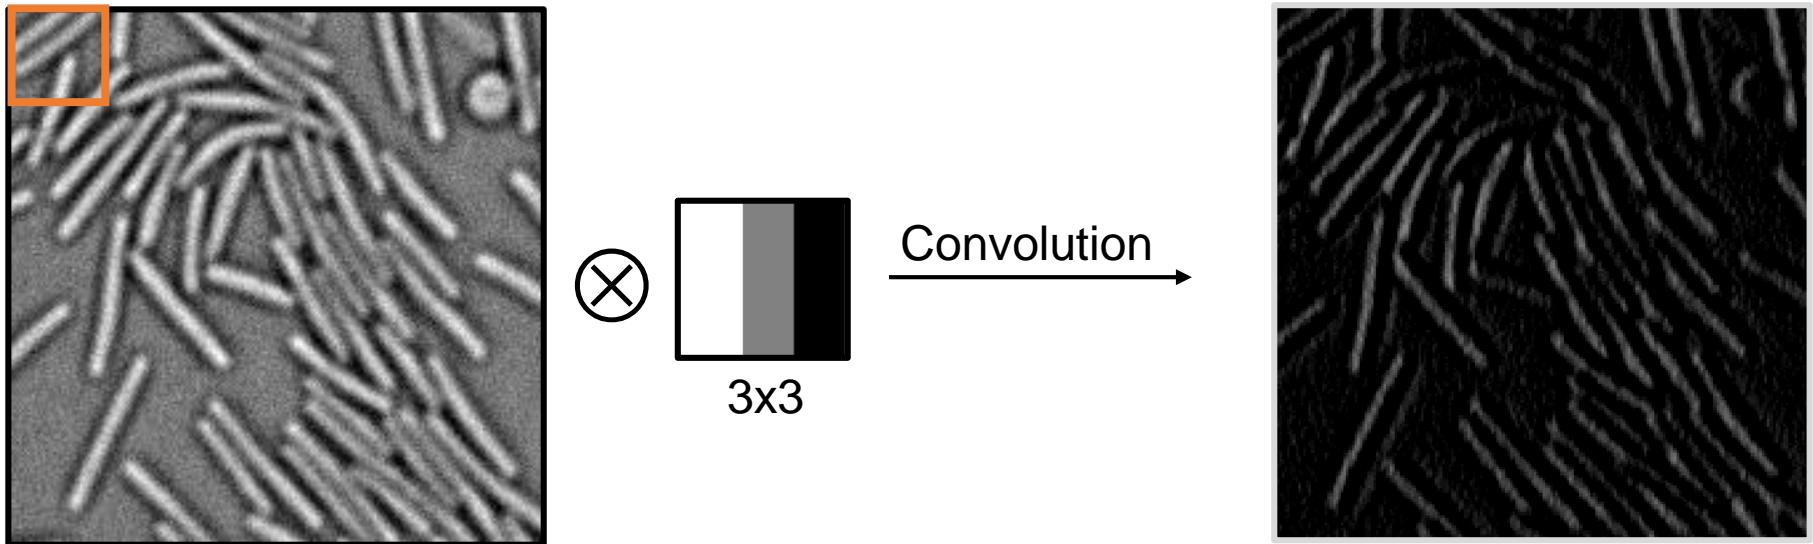
Densely connected networks are not suited for image analysis:

- Too many parameters, even for small networks
- Loses the local information around each pixel

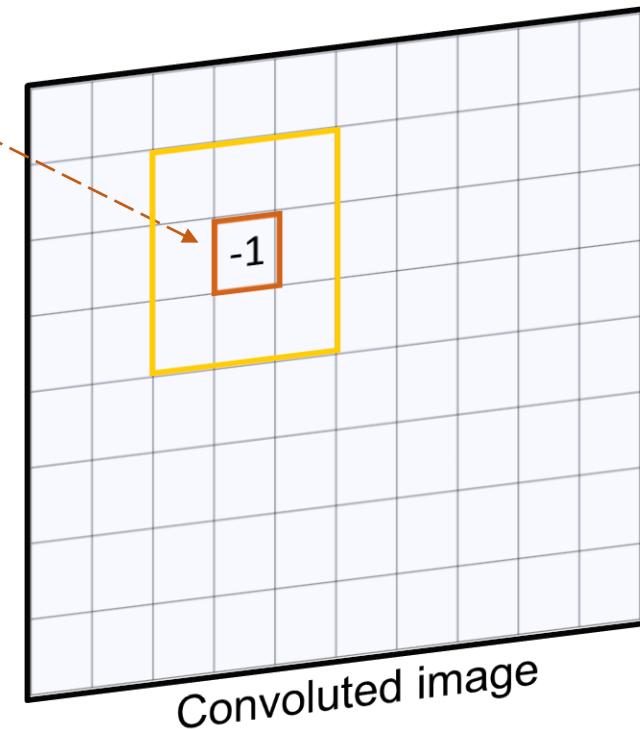
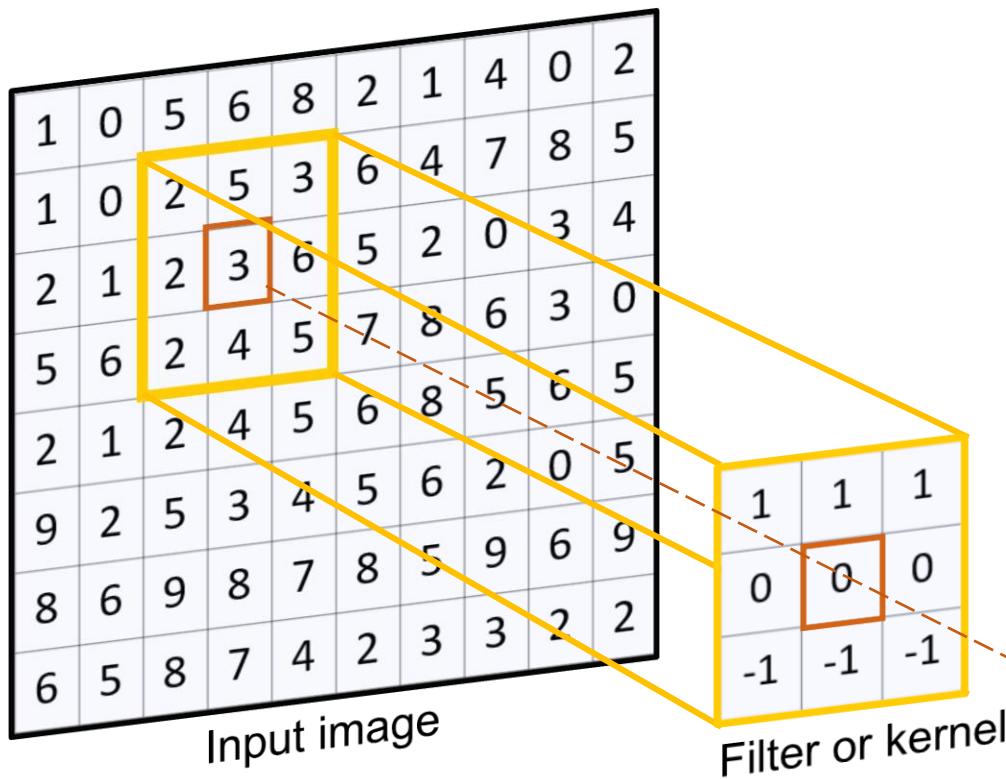


Learn local representations  
with convolutional network.

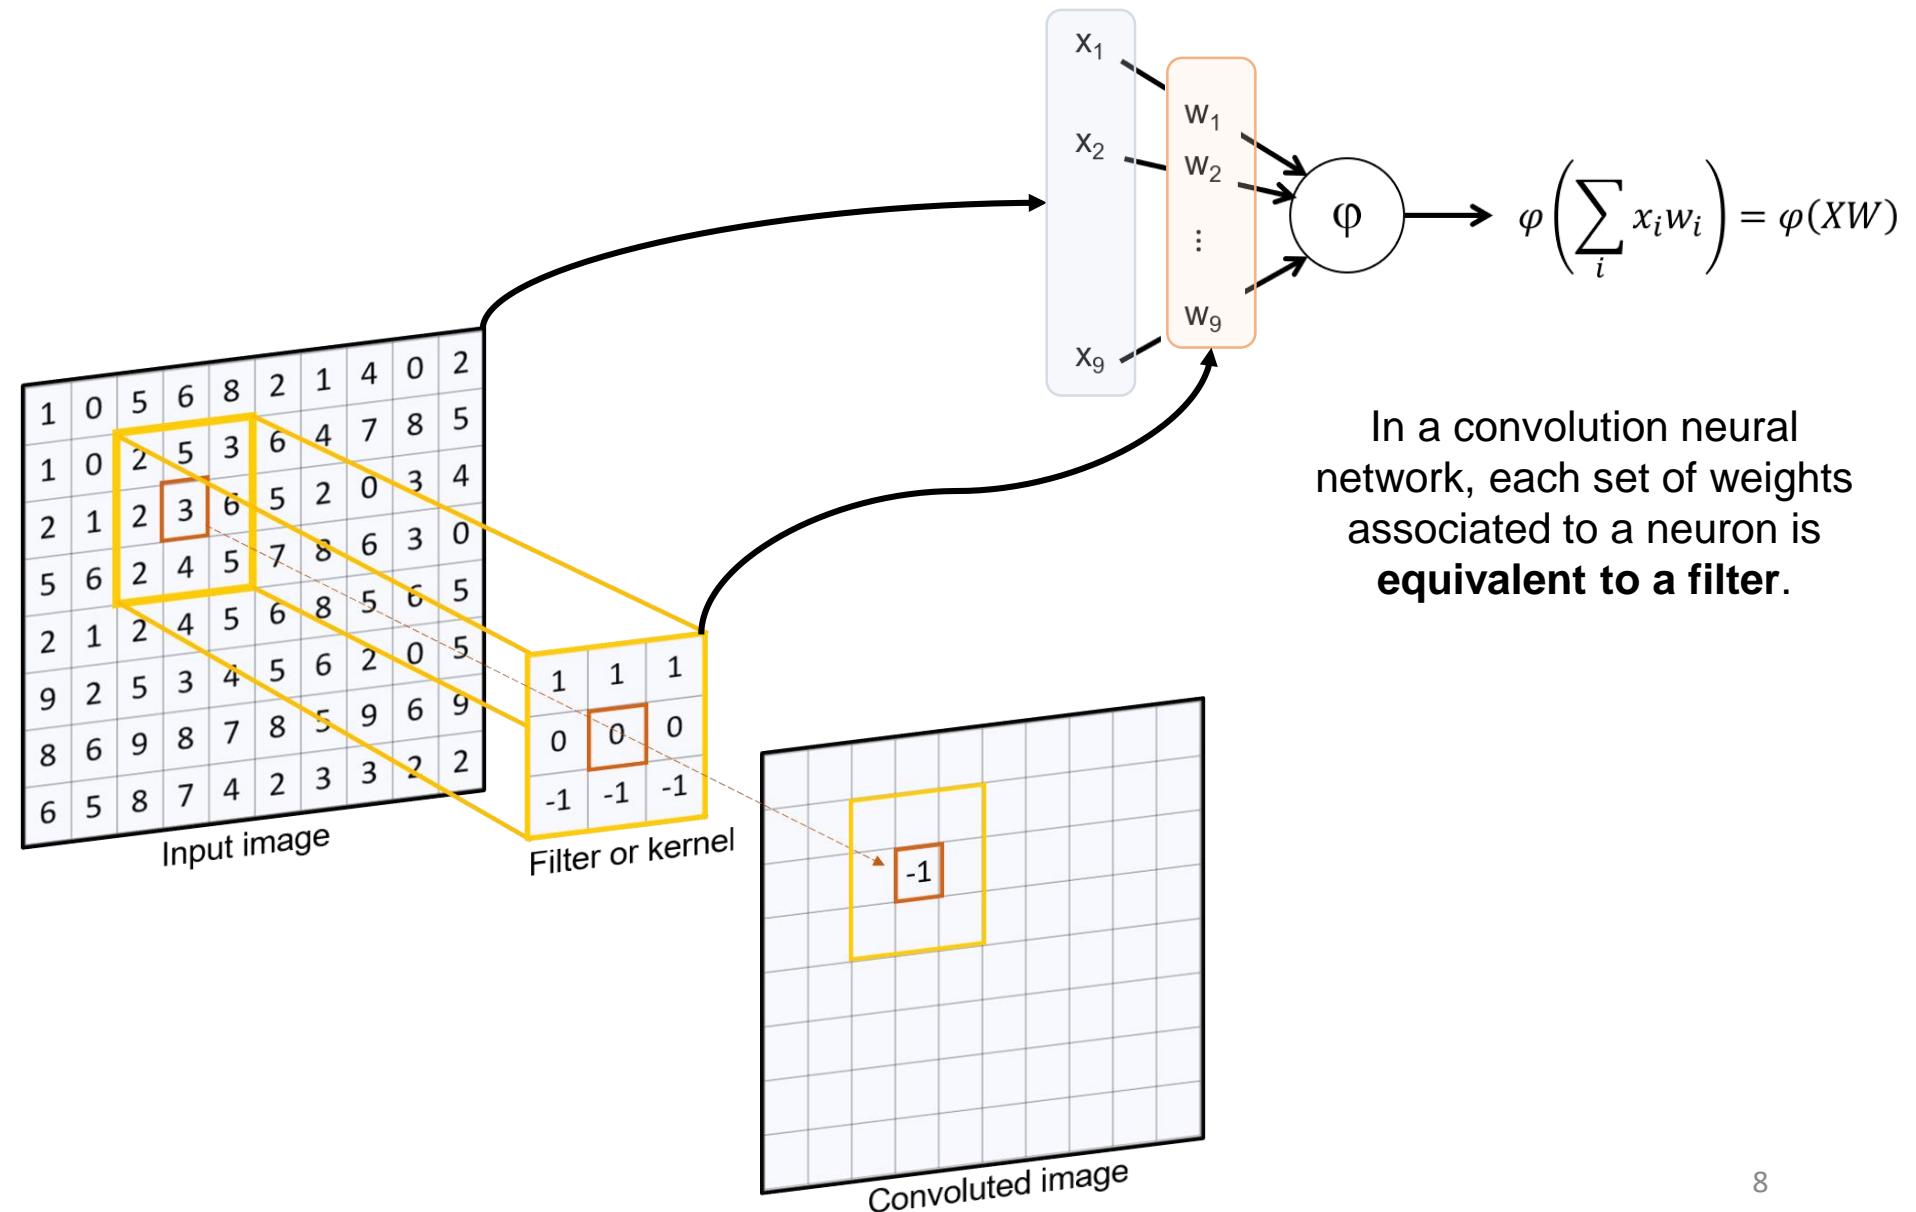
# Convolutional neural network



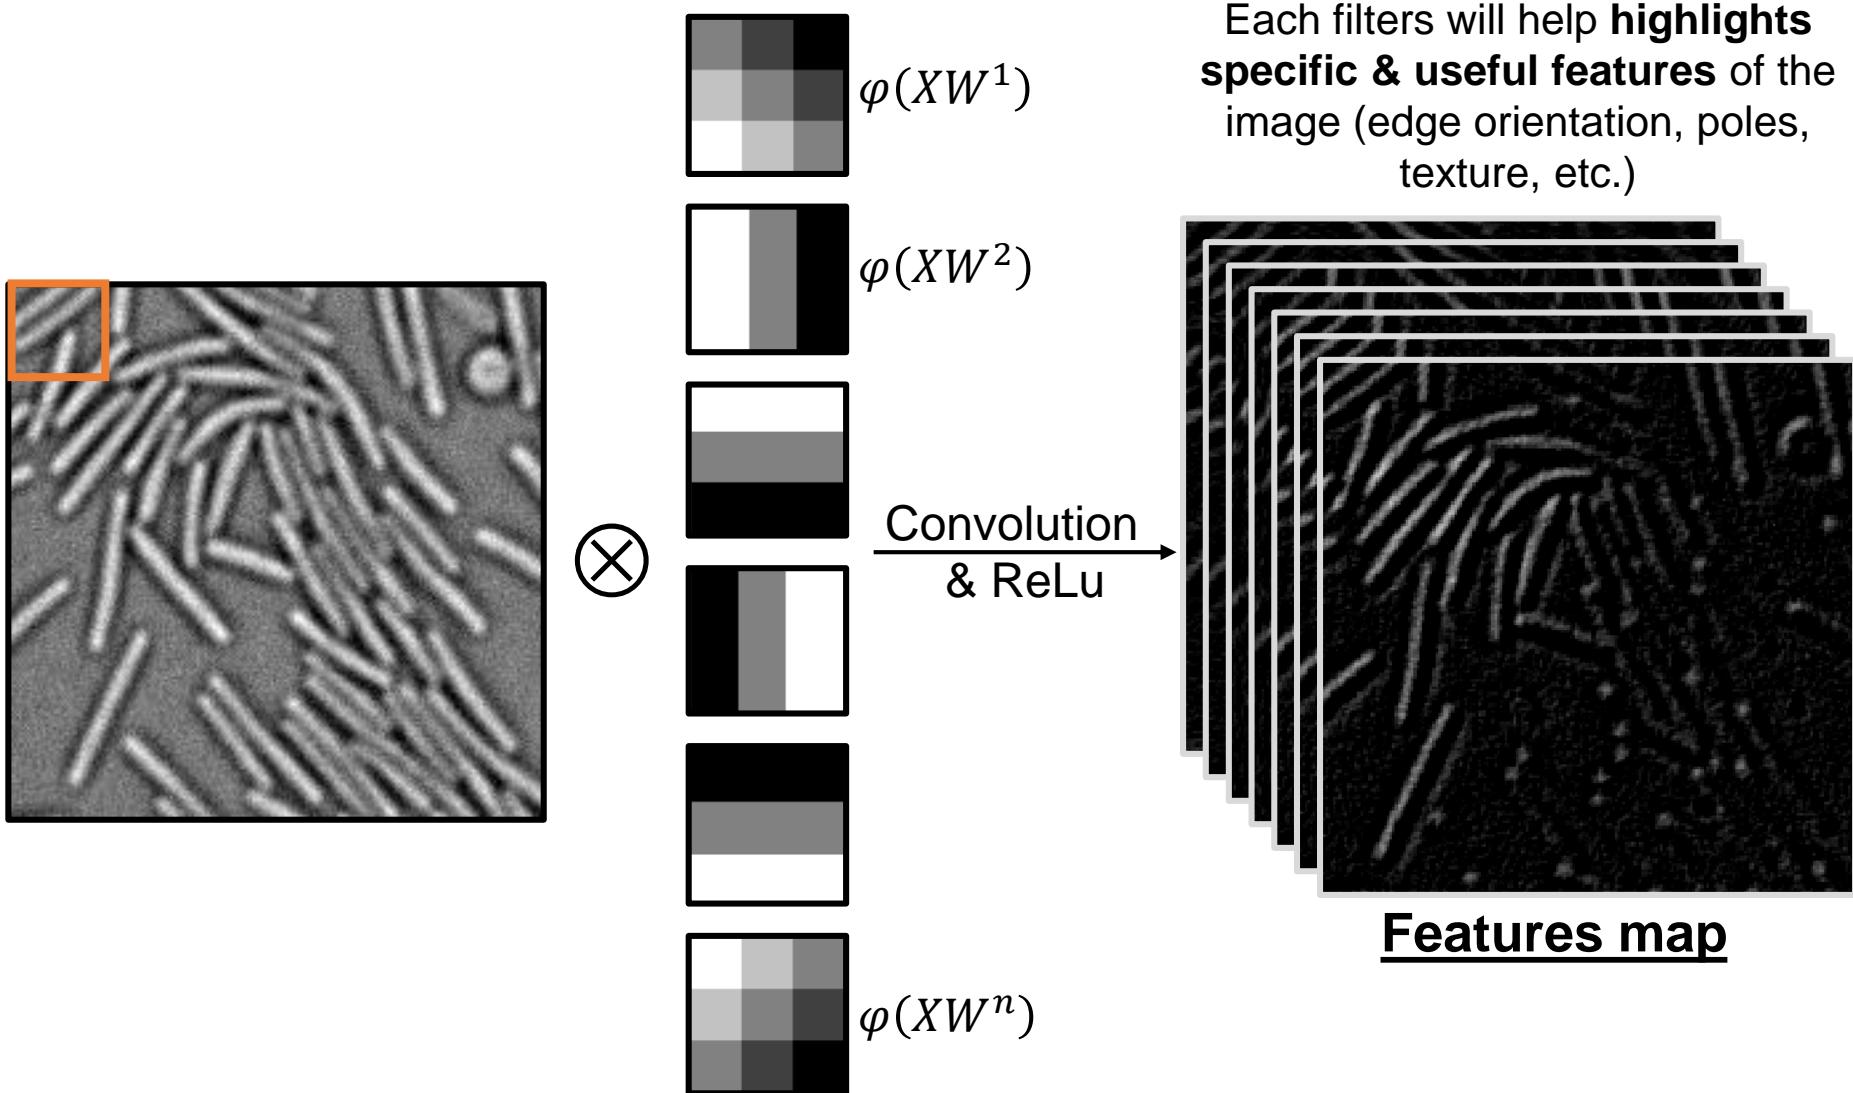
# Convolutional operation ...



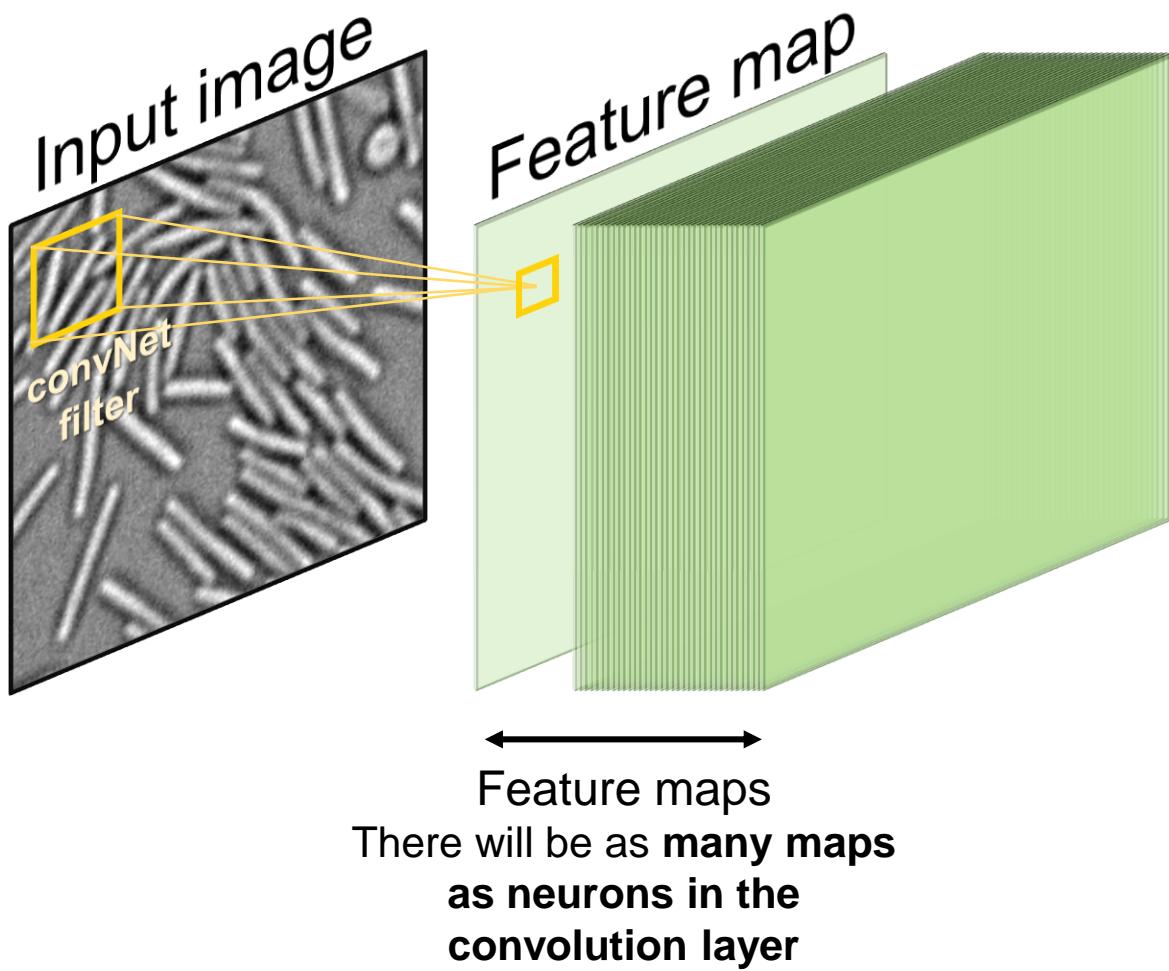
# Convolutional operation and single neuron operation



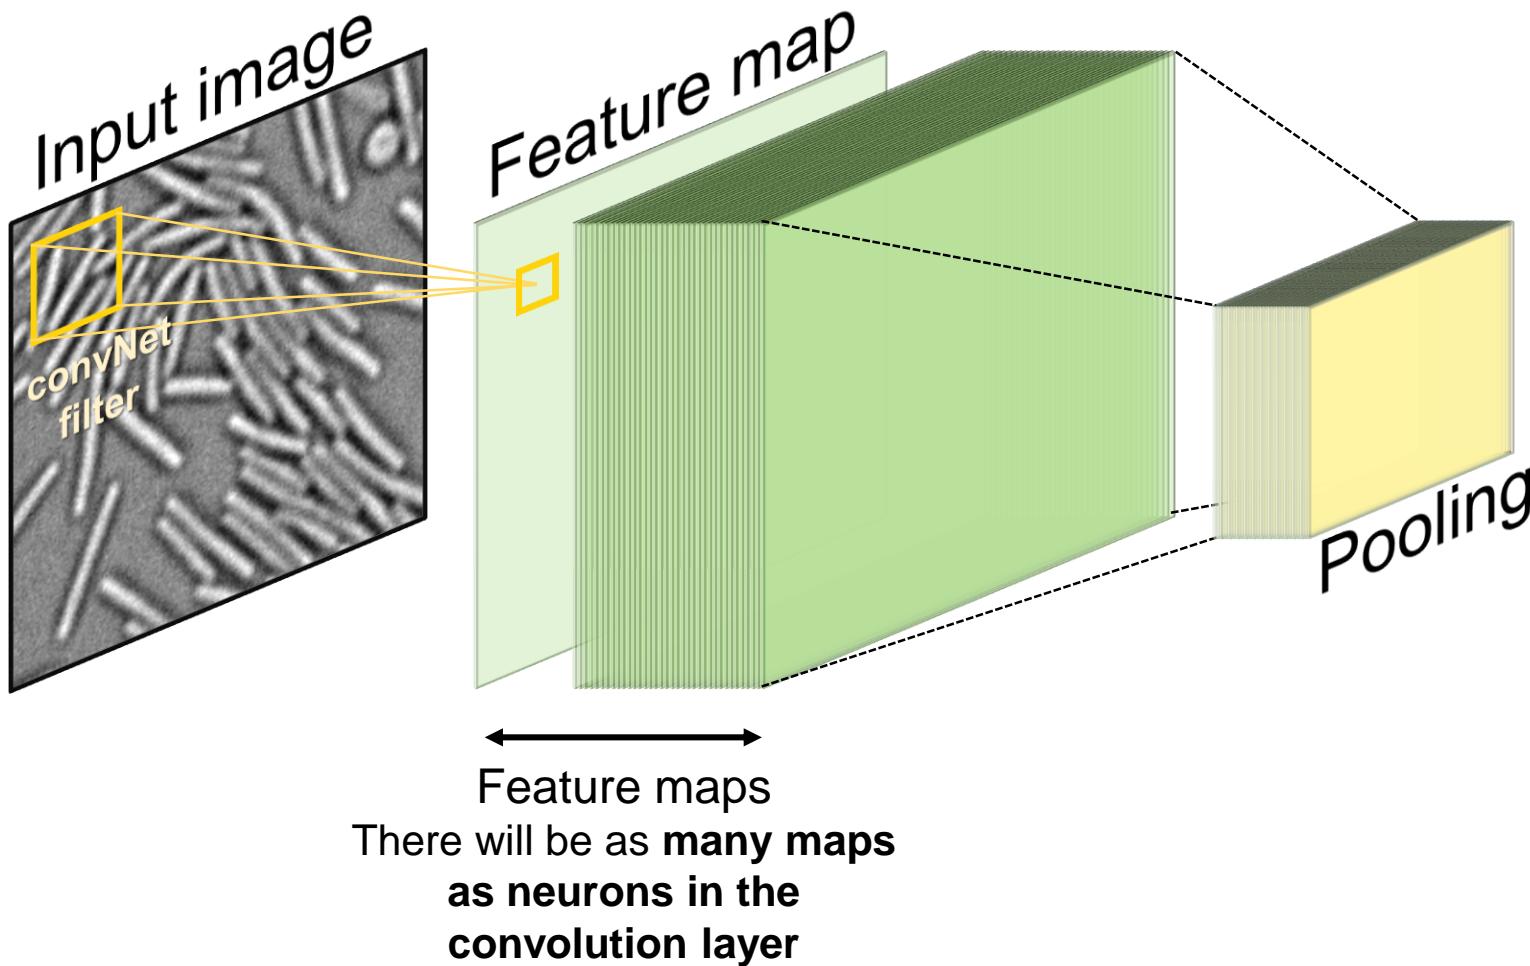
# Convolutional neural network



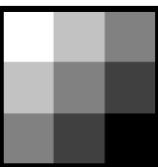
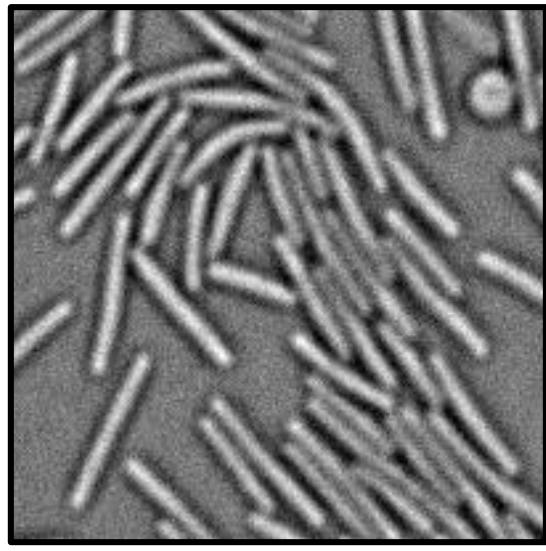
# Convolutional neural network



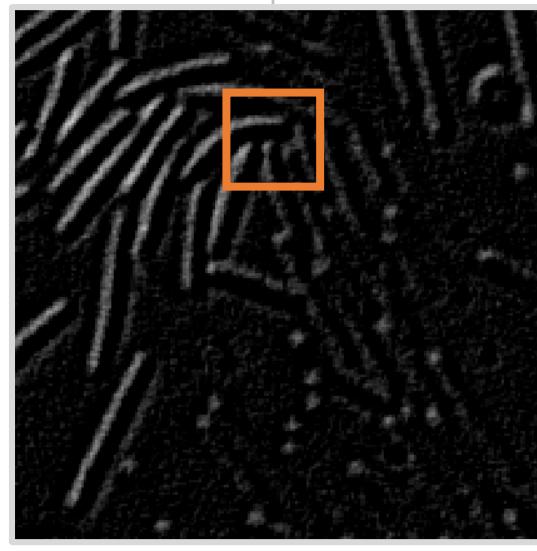
# Convolutional neural network



# Goal of the max-pooling layer ?



=



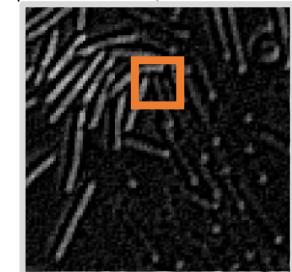
200 x 200 pixels

200 x 200 pixels

1	4	5	2
0	1	6	2
7	0	0	4
1	5	2	2

Max pooling  
2x2

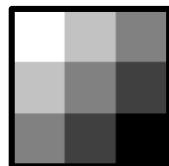
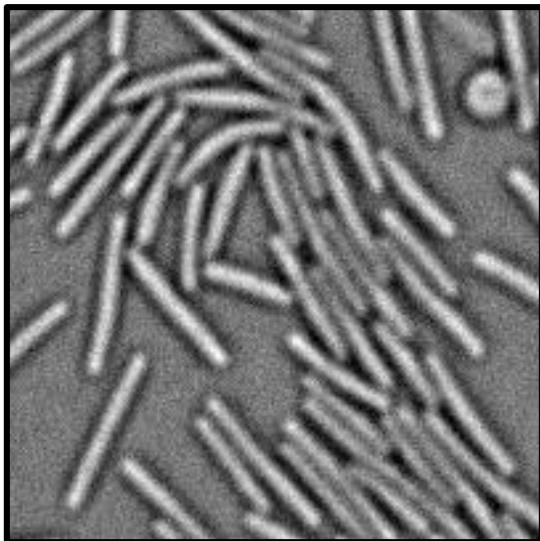
4	6
7	4



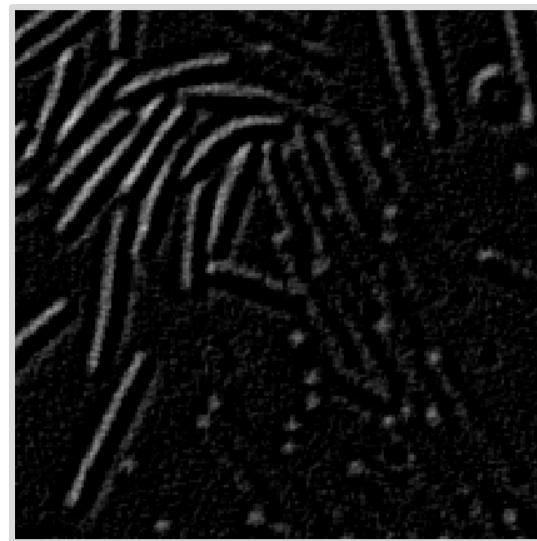
100 x 100  
pixels

# Goal of the max-pooling layer ?

1. **Reduce the spatial resolution** of the feature maps while keeping only the most relevant information
2. **Lowering memory and computing requirements**
3. Create **translation invariance**



=



200 x 200 pixels

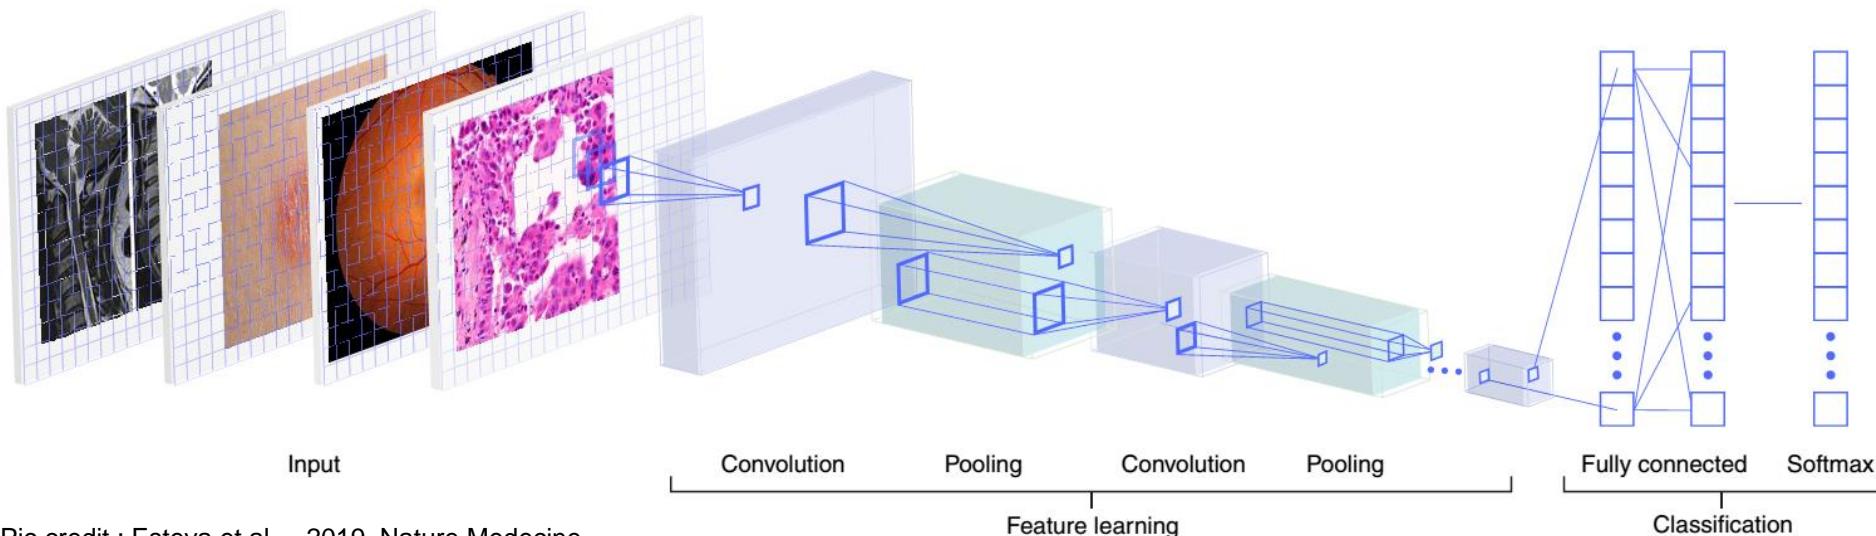
200 x 200 pixels

100 x 100  
pixels

# Convolutional network architecture

Classifier based on convNet are always built the same way:

1. A first part of the network is **using convolution layers to learn the features**
  - **Convolution layers → features extraction**
  - **Pooling → downsampling**
2. A second part is then classifying those features using **densely connected layers** in order to predict the right output.
  - Lots of parameters → **heavy on the memory**
  - Image input size is fixed → **not flexible**



# Outline

- I. Introduction**
- II. Starting with Deep Learning and Artificial Neuron Network**
- III. Deep Learning for image analysis**
  - i. Introduction to convolutional network
  - ii. Application : Digit recognition**
  - iii. Practical case 1 : image classification
  - iv. Practical case 2 : 2D regression of bead positions
  - v. Practical case 3 : Semantic segmentation of bacteria

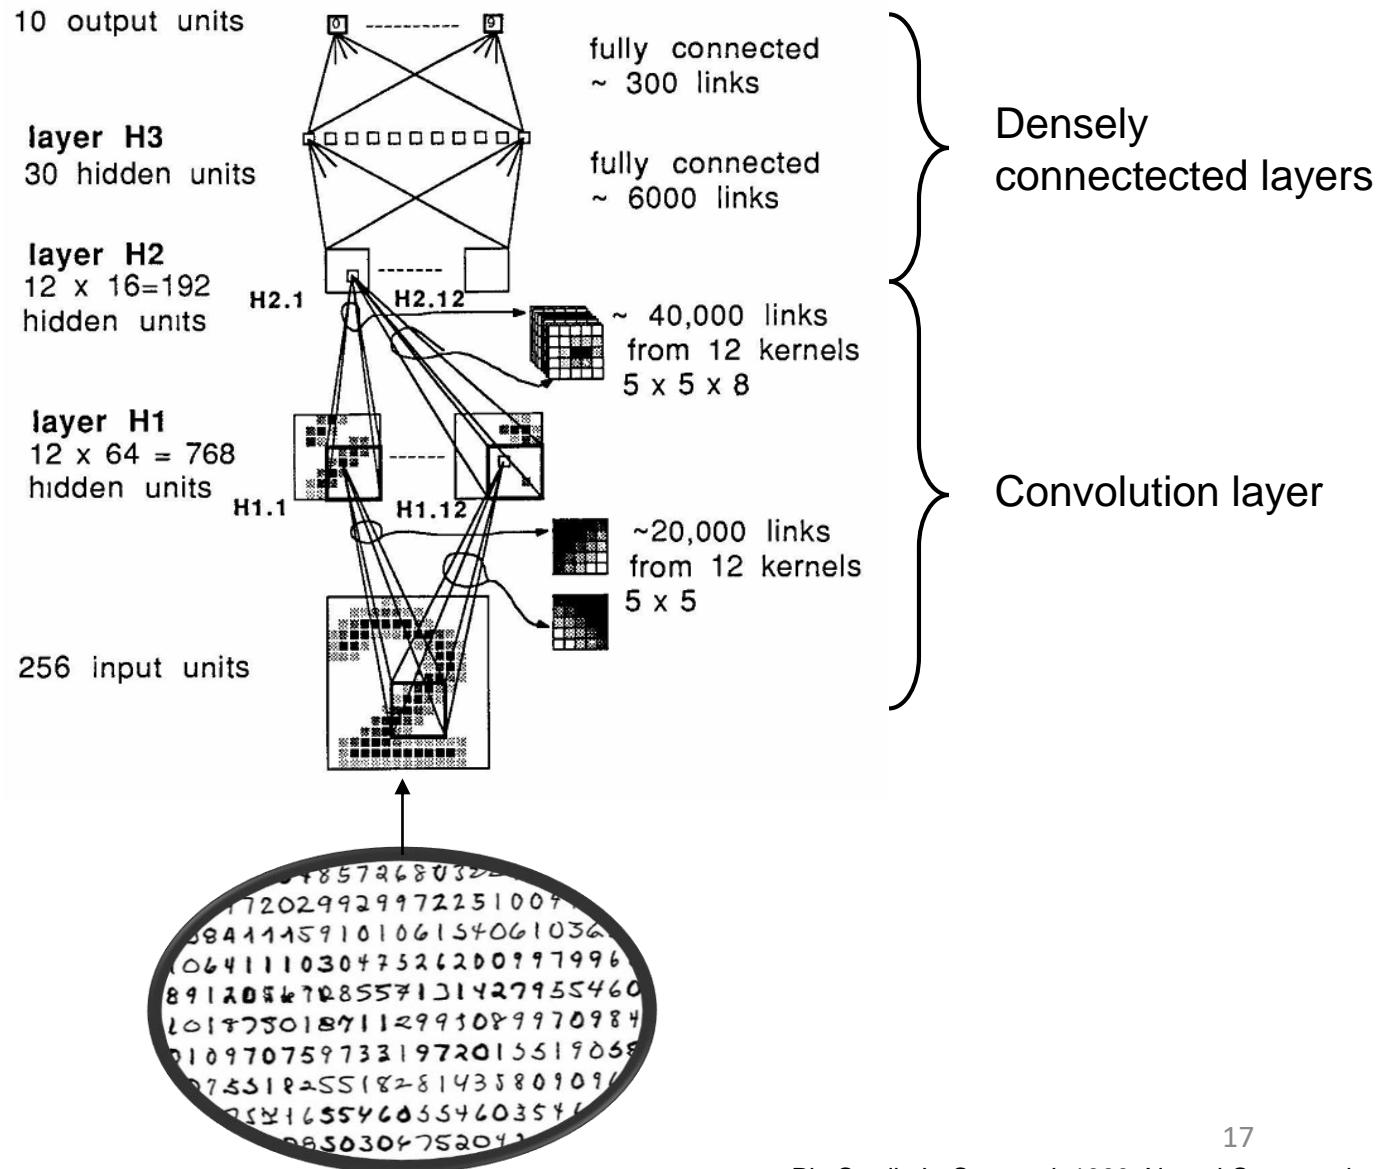
# Classify hand-written digits

## Example n°1 : MNIST\_convNet

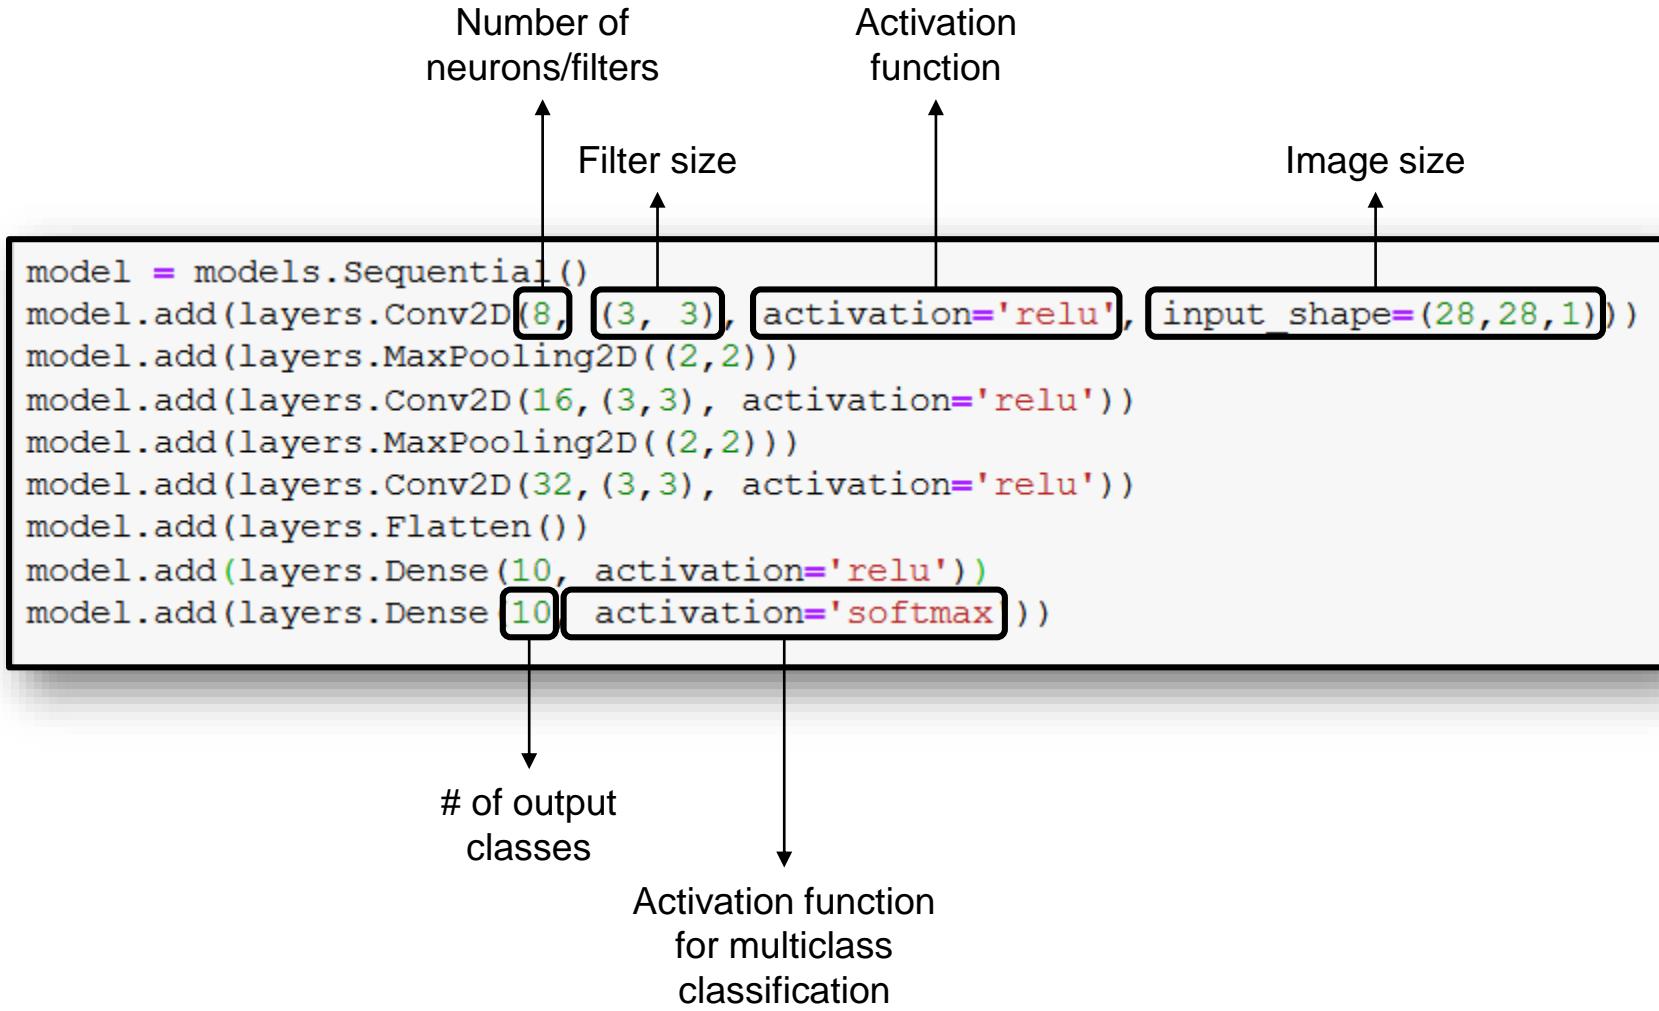
1. Understand and build your own ConvNet
2. Test this architecture for image classification



# Hand-written zipcode classification - 1989



# Hand-written digit classification with a ConvNet



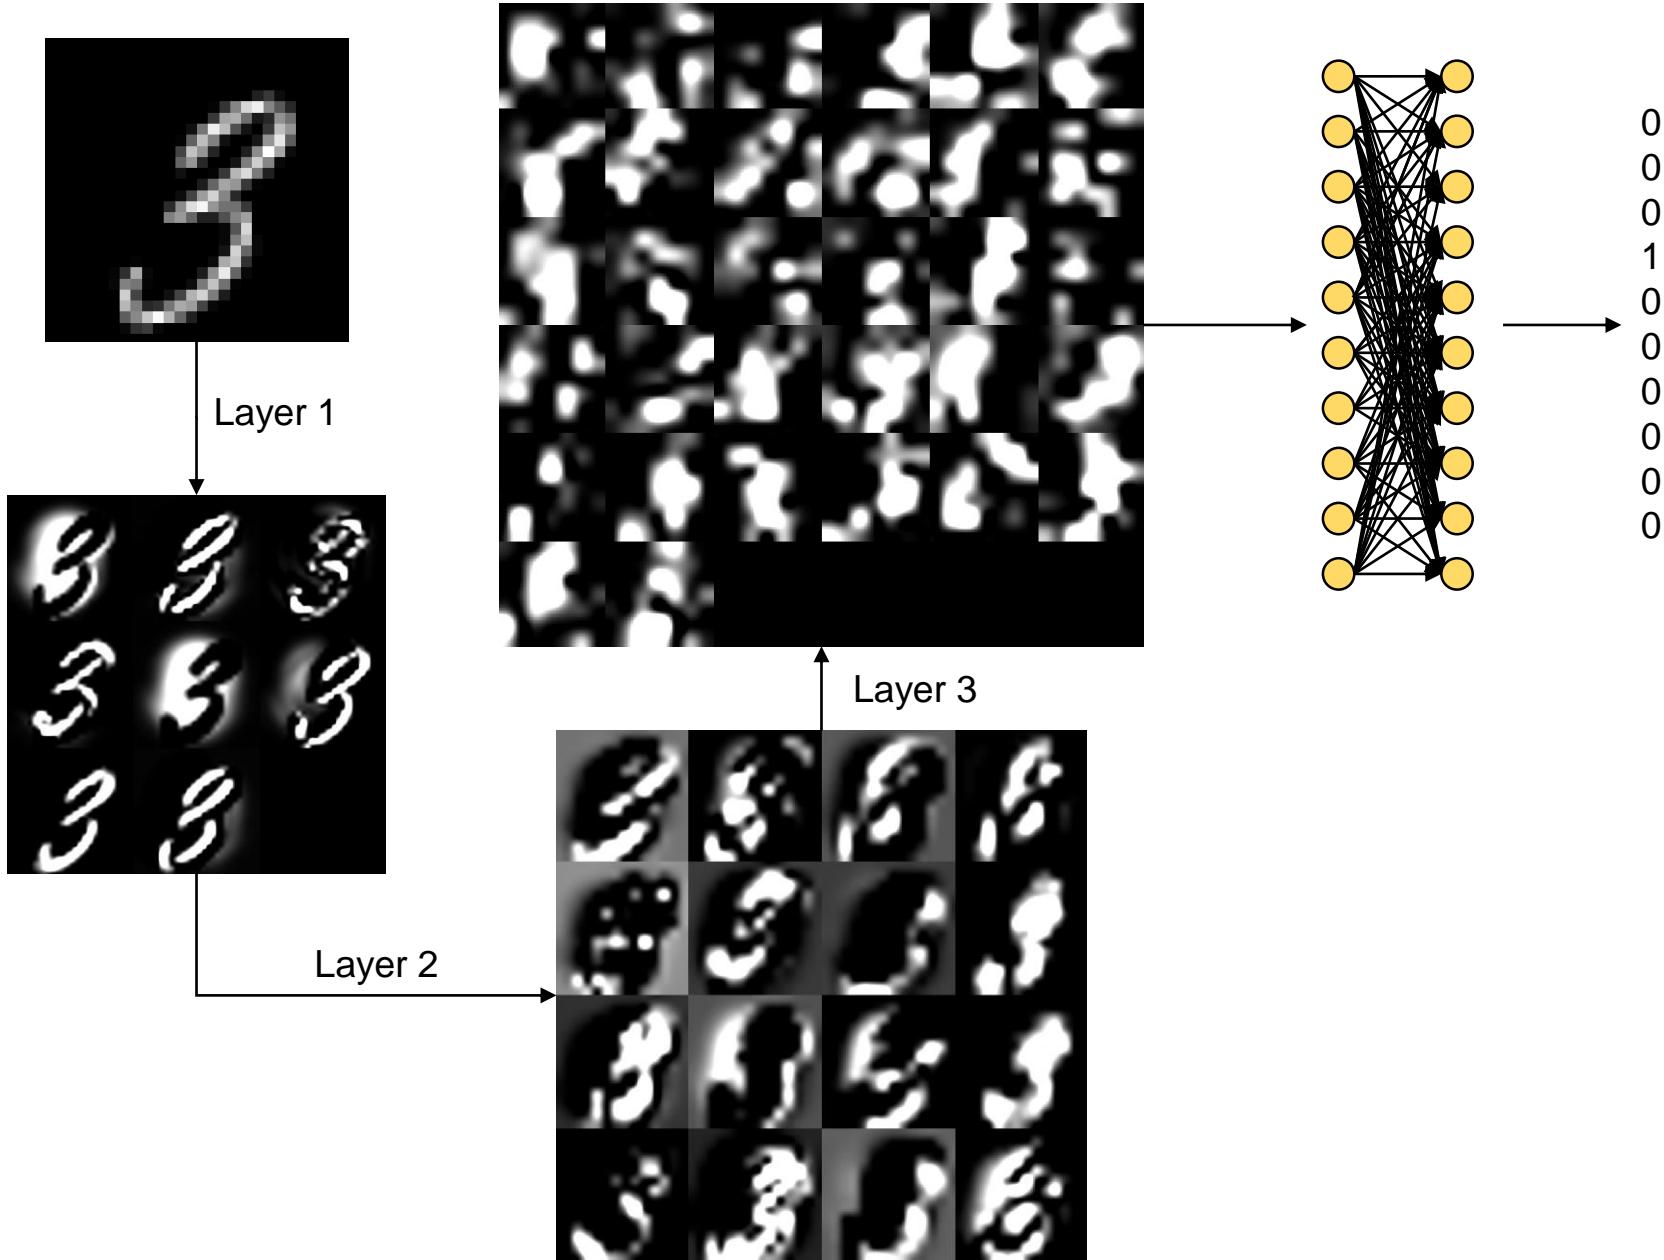
# Hand-written digit classification with a ConvNet

```
model = models.Sequential()
model.add(layers.Conv2D(8, (3, 3), activation='relu', input_shape=(28,28,1)))
model.add(layers.MaxPooling2D((2,2)))
model.add(layers.Conv2D(16, (3,3), activation='relu'))
model.add(layers.MaxPooling2D((2,2)))
model.add(layers.Conv2D(32, (3,3), activation='relu'))
model.add(layers.Flatten())
model.add(layers.Dense(10, activation='relu'))
model.add(layers.Dense(10, activation='softmax'))
```

Layer (type)	Output Shape	Param #
conv2d_11 (Conv2D)	(None, 26, 26, 8)	80
max_pooling2d_10 (MaxPooling)	(None, 13, 13, 8)	0
conv2d_12 (Conv2D)	(None, 11, 11, 16)	1168
max_pooling2d_11 (MaxPooling)	(None, 5, 5, 16)	0
conv2d_13 (Conv2D)	(None, 3, 3, 32)	4640
flatten_5 (Flatten)	(None, 288)	0
dense_8 (Dense)	(None, 10)	2890
dense_9 (Dense)	(None, 10)	110
<hr/>		
Total params: 8,888		
Trainable params: 8,888		
Non-trainable params: 0		

>30% of the parameters

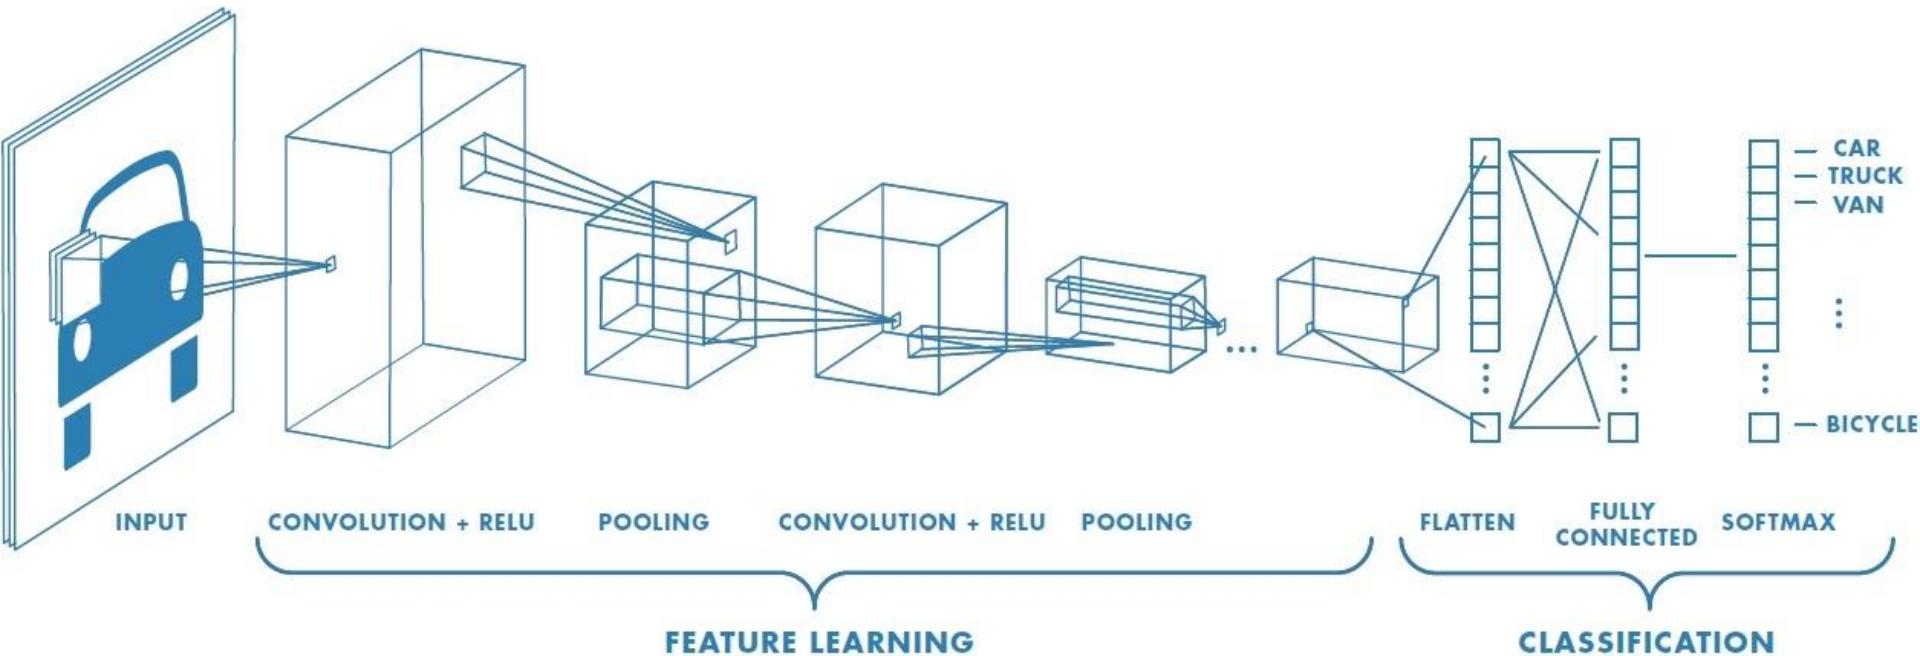
# What are the features?



# What did we learn?

How to build & train a simple classifier based on a convNet

```
model.add(layers.Conv2D(8, (3, 3), activation='relu', input_shape=(28,28,1)))  
  
model.add(layers.MaxPooling2D((2,2)))  
  
model.add(layers.Flatten())  
  
model.add(layers.Dense(10, activation='relu'))
```



# Outline

- I. Introduction**
- II. Starting with Deep Learning and Artificial Neuron Network**
- III. Deep Learning for image analysis**
  - i. Introduction to convolutional network
  - ii. Application : Digit recognition
  - iii. Practical case 1 : image classification**
  - iv. Practical case 2 : 2D regression of bead positions
  - v. Practical case 3 : Semantic segmentation of bacteria

# Deep Learning applied to image classification



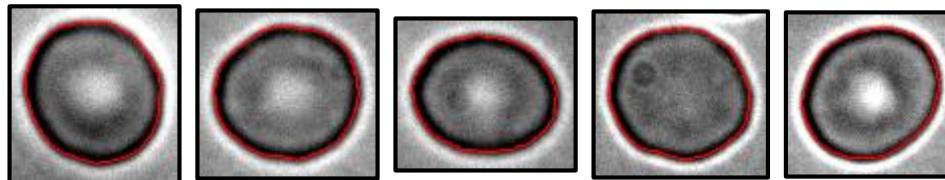
Viviana Claveria



Manouk Abkarian



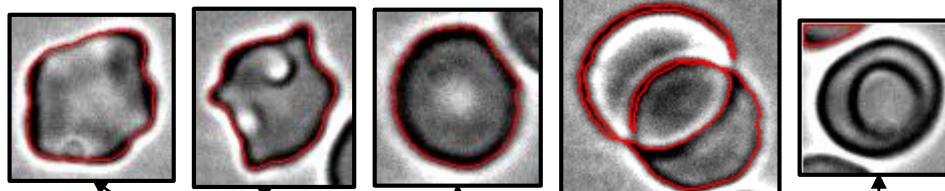
Good images



In-focus image of RBC with a properly defined contour (red)



Bad images



Poor quality images that need to be discarded from the analysis

Sick cells

Out-of-focus

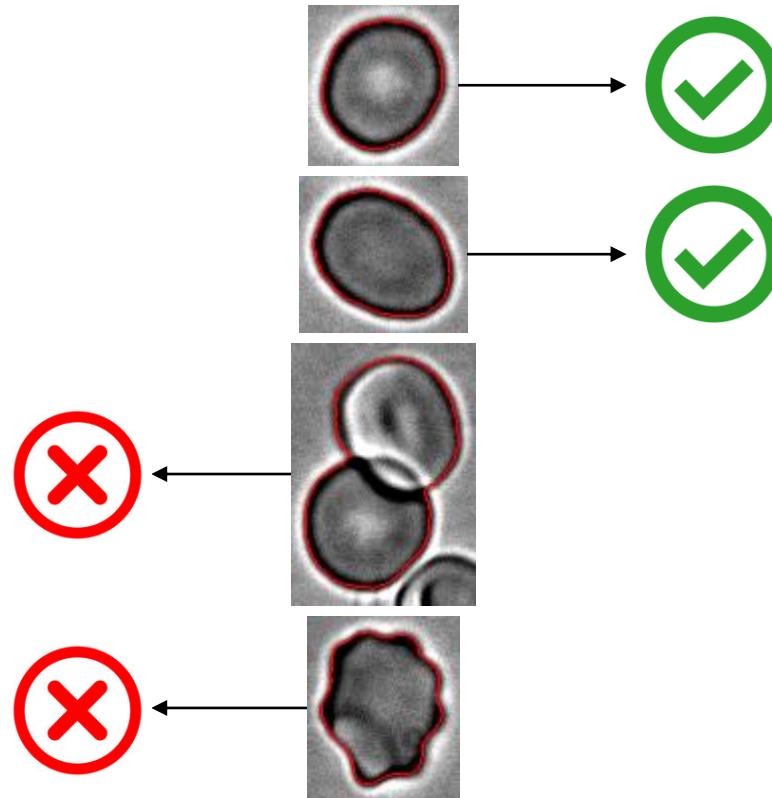
Overlapping

Failed contour

# Classify red blood cell (RBC) images

## Example n°2 : RBC\_image\_classification\_no\_image\_augmentation

1. Build your own ConvNet classifier to solve a biological problem
2. Understand how to detect overfitting and which strategies to avoid it
3. How already-trained network can be used for classification



# One possible solution ...

Type of data : png RGB images resize to 50x50 and normalized

Size of the training set : 956 good images / 2000 bad images

Size of the validation/testing sets : 319 good images / 667 bad images

```
model = models.Sequential()
model.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(50, 50, 3)))
model.add(layers.MaxPooling2D((2,2)))
model.add(layers.Conv2D(64, (3,3), activation='relu'))
model.add(layers.MaxPooling2D((2,2)))
model.add(layers.Conv2D(128, (3,3), activation='relu'))
model.add(layers.MaxPooling2D((2,2)))
model.add(layers.Conv2D(128, (3,3), activation='relu'))
model.add(layers.Flatten())
model.add(layers.Dense(256, activation='relu'))
model.add(layers.Dense(128, activation='relu'))
model.add(layers.Dense(1, activation='sigmoid'))

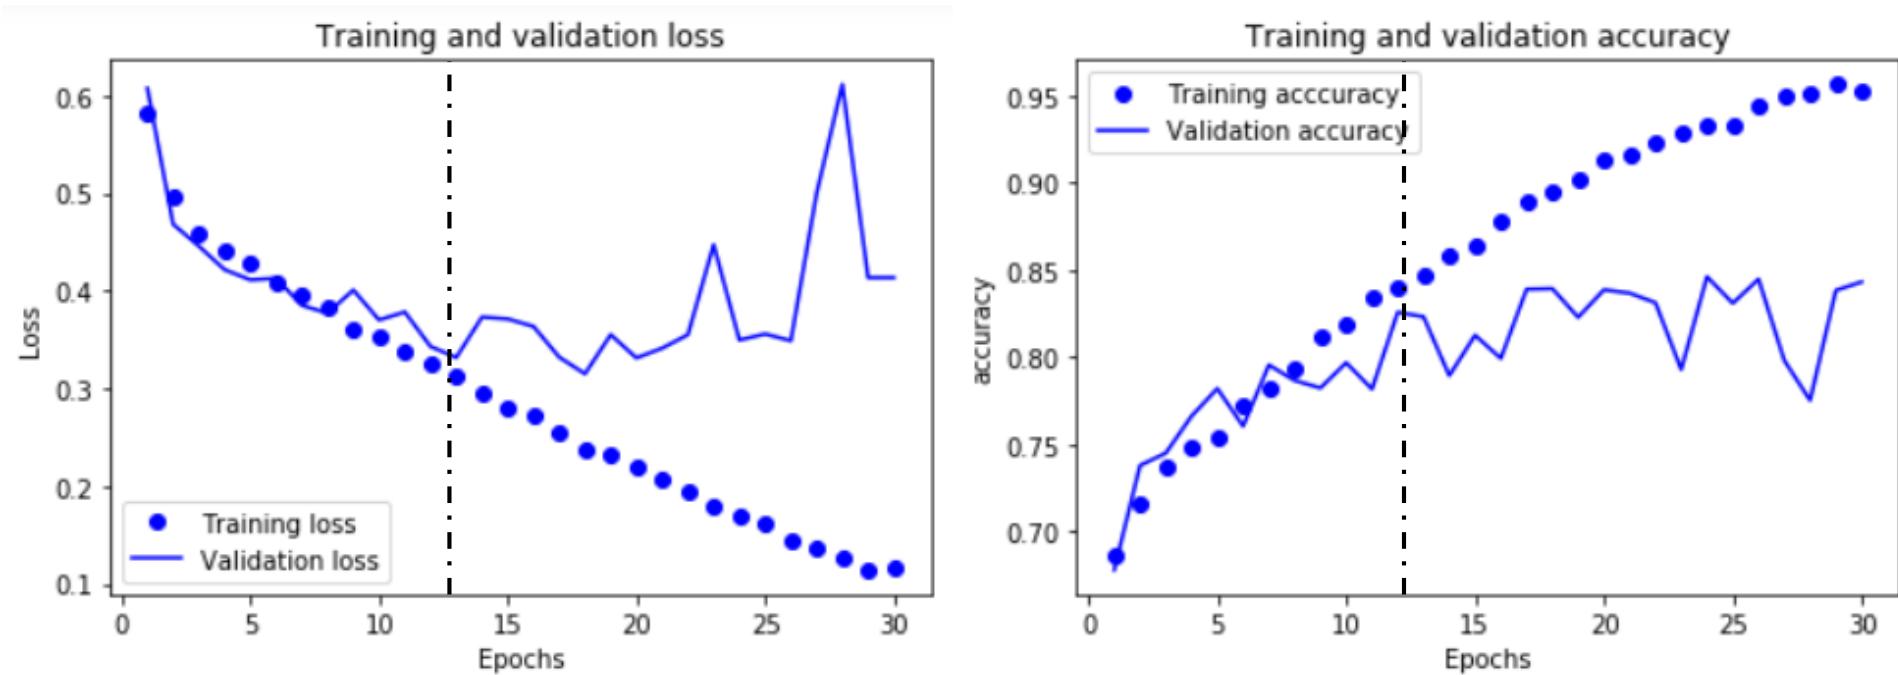
model.compile(optimizer = optimizers.RMSprop(lr=1e-4),
              loss='binary_crossentropy',
              metrics=['accuracy'])

model.summary()
```

Number of trainable parameters : 405,185



# ... but overfitting

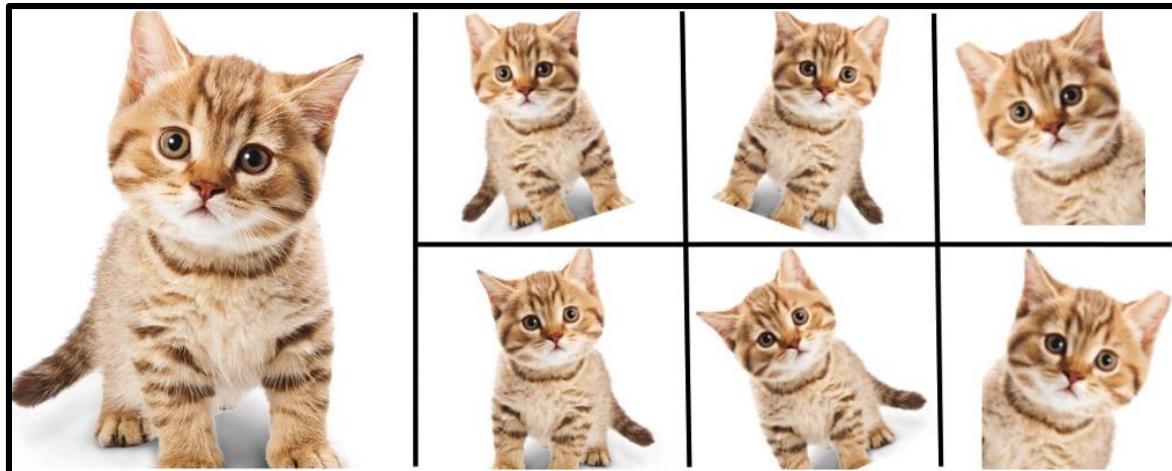


Using the validation set, we observe that the training and validation loss & accuracy are diverging after epoch #12. **The network is no longer learning new useful features.**

- Overfitting
- Validation loss & accuracy are noisy
- The global accuracy of the network is ~84%

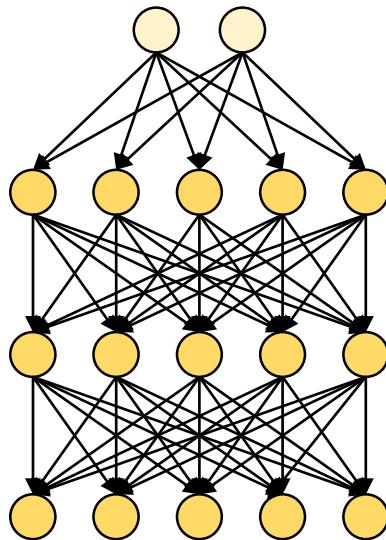
# How to avoid overfitting?

- **Reduce the size of the network** ... but no magical formula to determine how many layers we need
- Increase the size of the training set:
  - Actually **increase the number of images**
  - Use **image augmentation** strategy

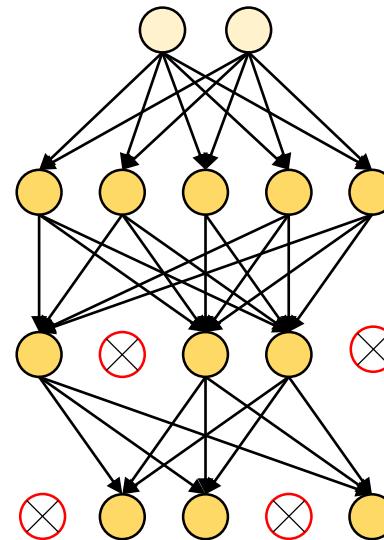


# How to avoid overfitting?

- **Reduce the size of the network** ... but no magical formula to determine how many layers we need
- Increase the size of the training set:
  - Actually **increase the number of images**
  - Use **image augmentation** strategy
- **Dropout**, randomly “turning-off” neurons of the network



No dropout



Dropout with 40% probability

Dropout is used to avoid co-adaptation of neurons → enforce the fact that neurons should **learn and work independently**.

# How to avoid overfitting?

- **Reduce the size of the network** ... but no magical formula to determine how many layers we need
- Increase the size of the training set:
  - Actually **increase the number of images**
  - Use **image augmentation** strategy
- **Dropout**, randomly “turning-off” neurons of the network
- **Weight regularization  $L_1$  &  $L_2$** , a strategy to force the weights to take only small values during the training:
  - $L_1$  – the cost added to the loss is proportional to the absolute weight value
  - $L_2$  – the cost is proportional to the square value of the weight

# Effect of image augmentation & dropout

Type of data : png RGB images resize to 85x85 and normalized

Size of the training set : 956 good images / 2000 bad images

Size of the validation/testing sets : 319 good images / 667 bad images

```
model = models.Sequential()
model.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(85, 85, 3)))
model.add(layers.MaxPooling2D((2,2)))
model.add(layers.Conv2D(64, (3,3), activation='relu'))
model.add(layers.MaxPooling2D((2,2)))
model.add(layers.Conv2D(128, (3,3), activation='relu'))
model.add(layers.MaxPooling2D((2,2)))
model.add(layers.Conv2D(128, (3,3), activation='relu'))
model.add(layers.Flatten())
model.add(layers.Dense(256, activation='relu'))
model.add(layers.Dropout(0.5))
model.add(layers.Dense(128, activation='relu'))
model.add(layers.Dropout(0.5))
model.add(layers.Dense(1, activation='sigmoid'))

model.compile(optimizer = optimizers.RMSprop(lr=1e-4),
              loss='binary_crossentropy',
              metrics=['accuracy'])

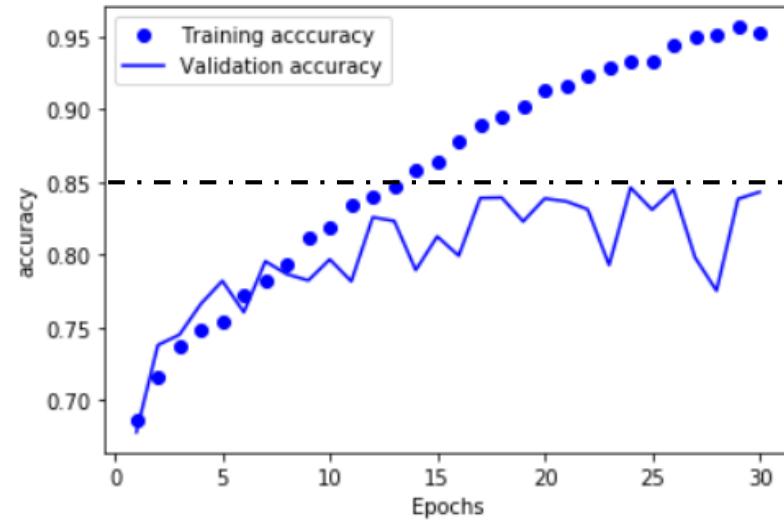
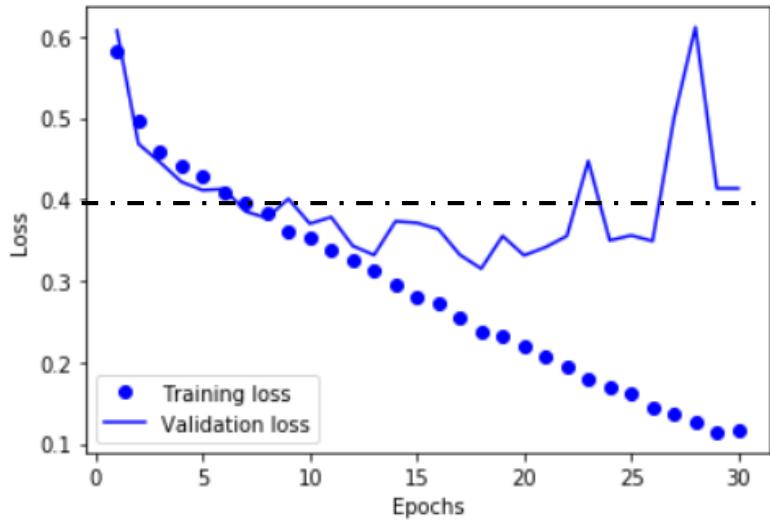
model.summary()
```



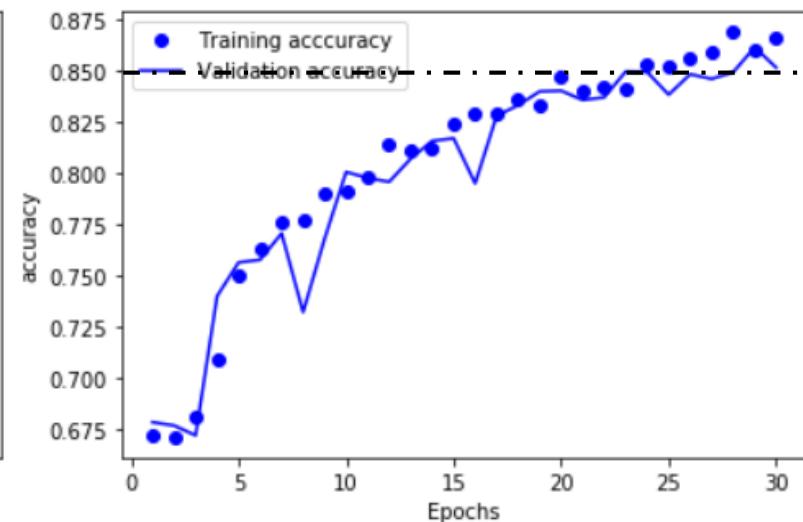
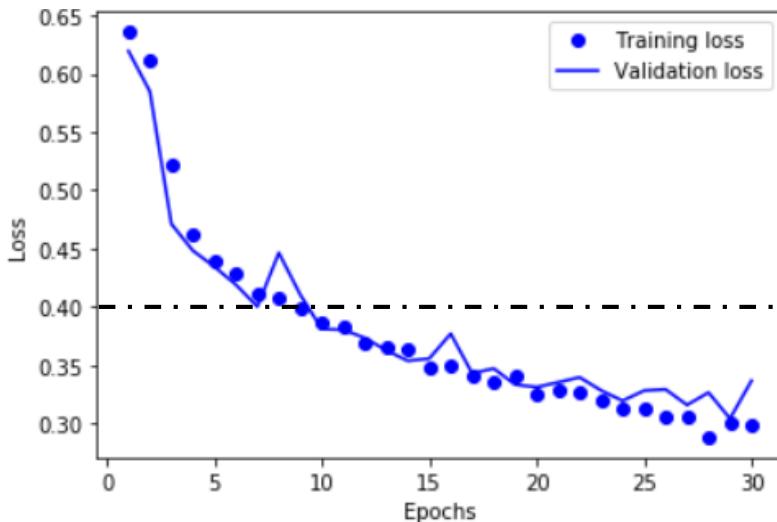
**Number of trainable parameters : 1,453,761**

# Better but the accuracy is still <90%

## Basic convNet



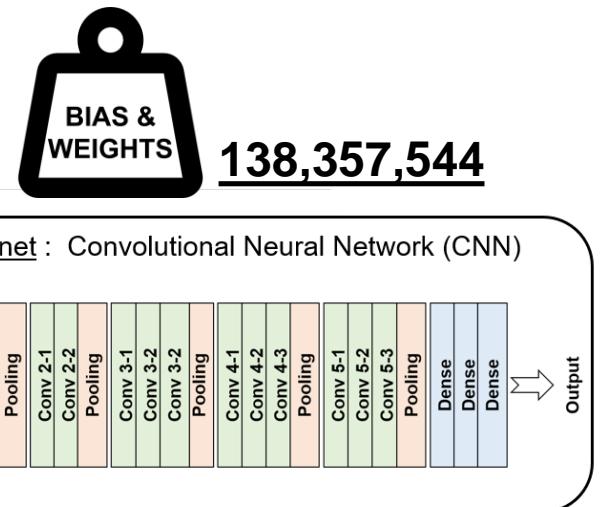
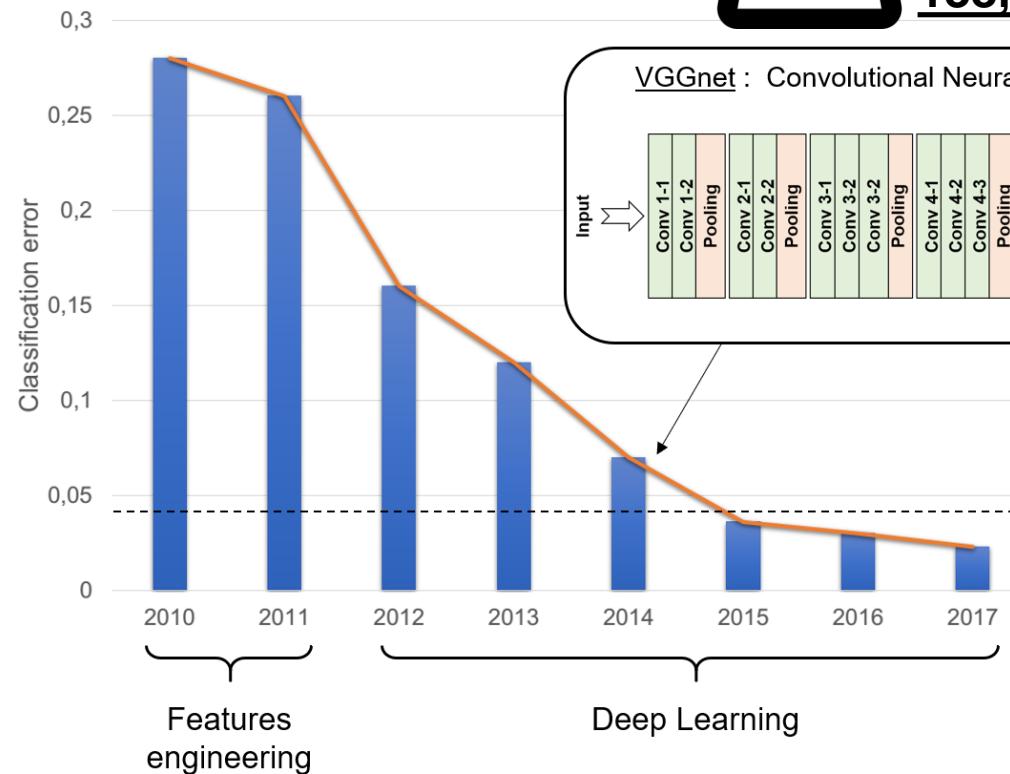
## ConvNet with image augmentation & dropout



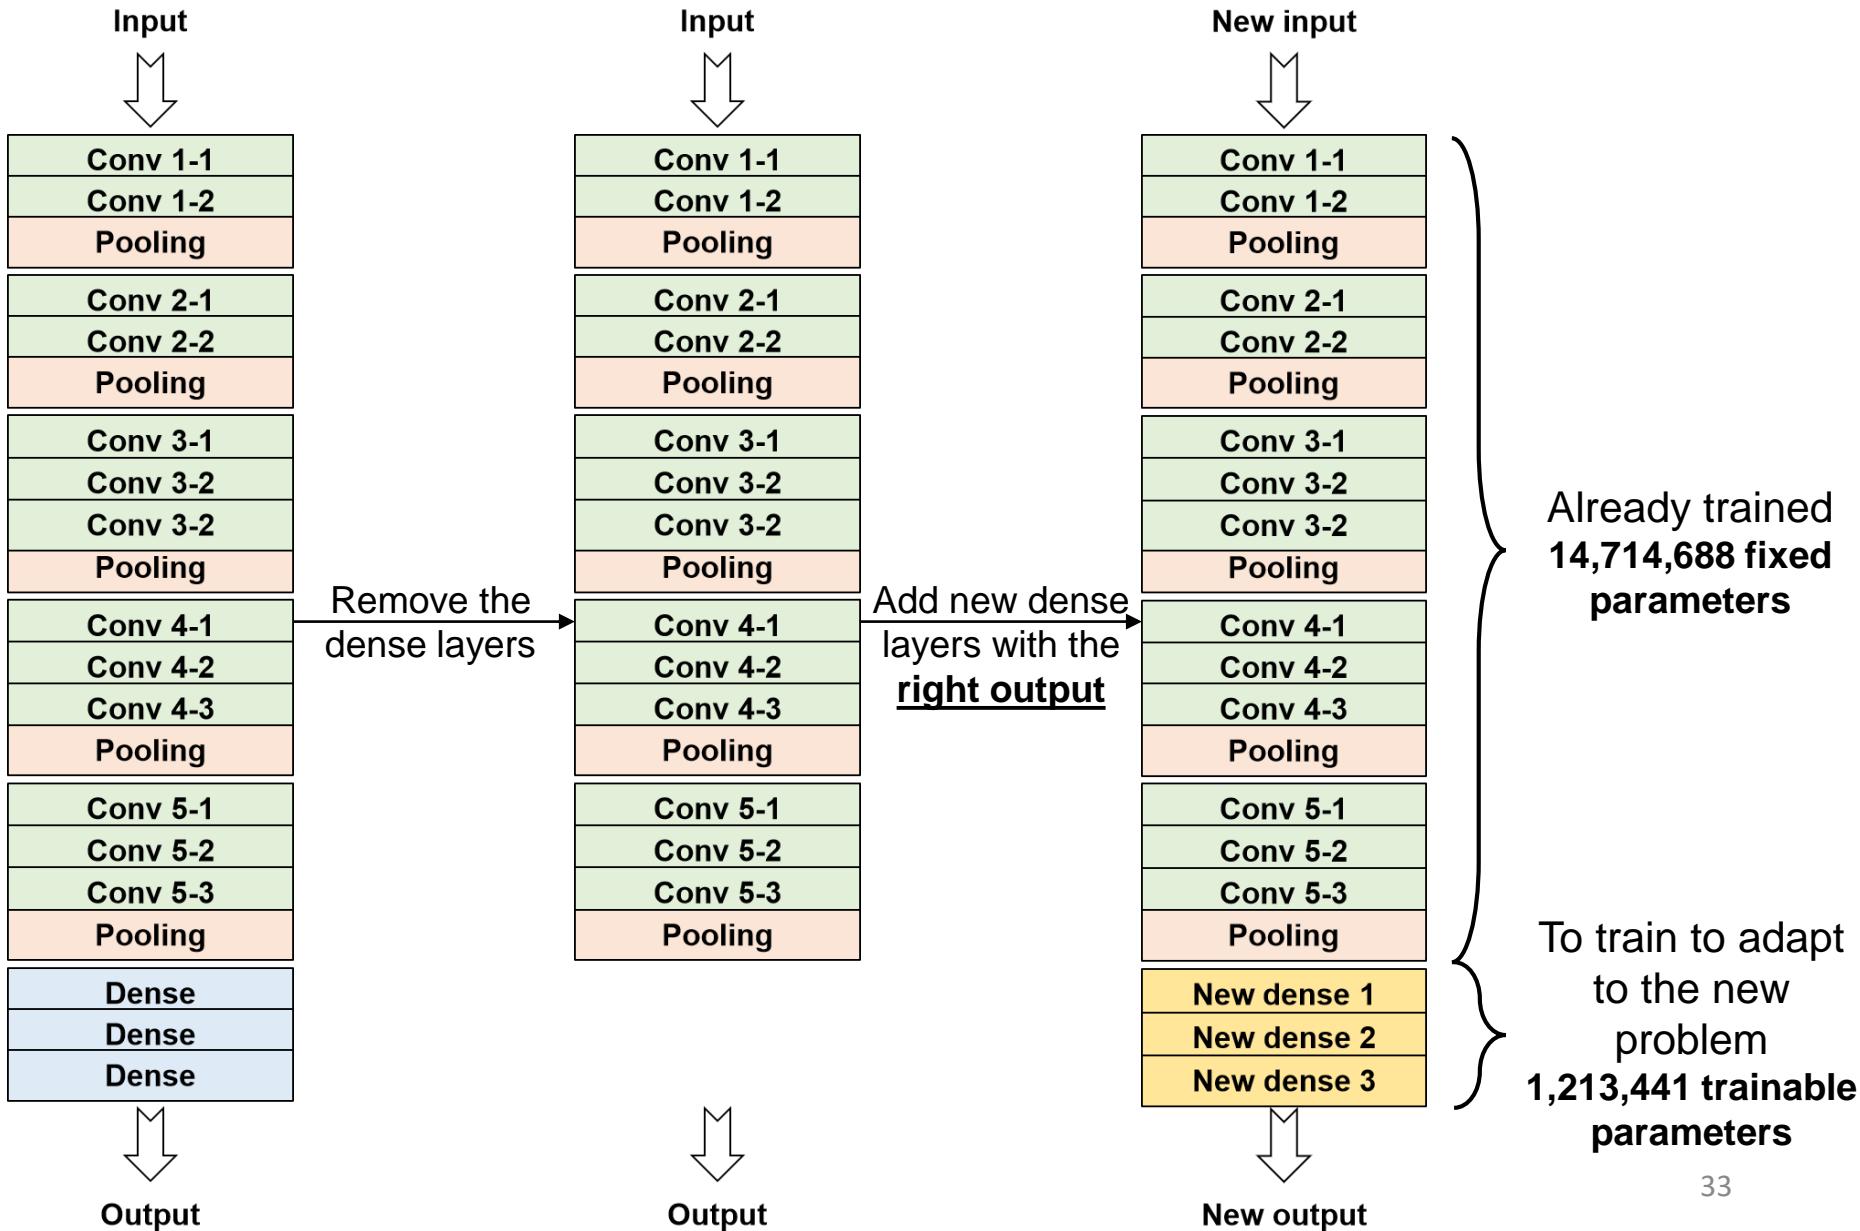
# Used a pre-trained convNet

Many networks used for the **Image Large Scale Visual Recognition Challenge** are available and already trained on million of images.

<https://keras.io/applications/>



# Used a pre-trained convNet : VGG16



# Used a pre-trained convNet : VGG16

```
from keras.applications import VGG16
conv_base = VGG16(weights='imagenet', include_top = False, input_shape=(85,85,3))

model = models.Sequential()
model.add(conv_base)
model.add(layers.Flatten())
model.add(layers.Dense(512, activation='relu'))
model.add(layers.Dropout(0.5))
model.add(layers.Dense(256, activation='relu'))
model.add(layers.Dropout(0.5))
model.add(layers.Dense(128, activation='relu'))
model.add(layers.Dense(1, activation='sigmoid'))

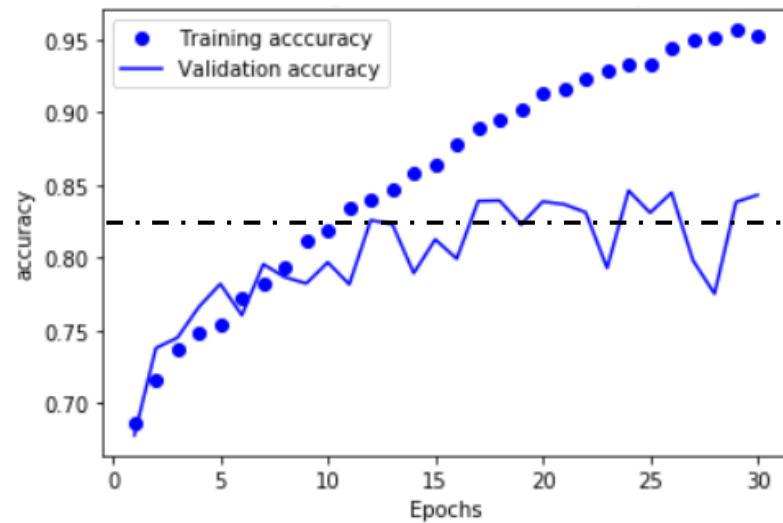
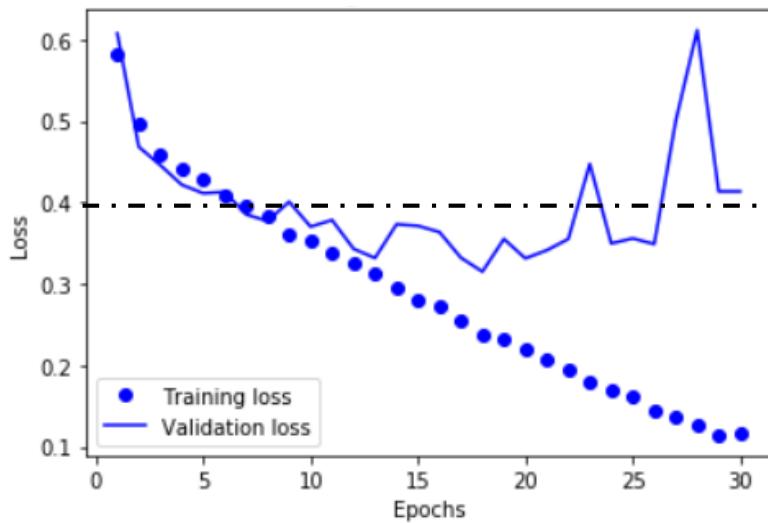
conv_base.trainable = False
```

Layer (type)	Output Shape	Param #
vgg16 (Model)	(None, 2, 2, 512)	14714688
flatten_6 (Flatten)	(None, 2048)	0
dropout_16 (Dropout)	(None, 2048)	0
dense_21 (Dense)	(None, 512)	1049088
dropout_17 (Dropout)	(None, 512)	0
dense_22 (Dense)	(None, 256)	131328
dropout_18 (Dropout)	(None, 256)	0
dense_23 (Dense)	(None, 128)	32896
dense_24 (Dense)	(None, 1)	129

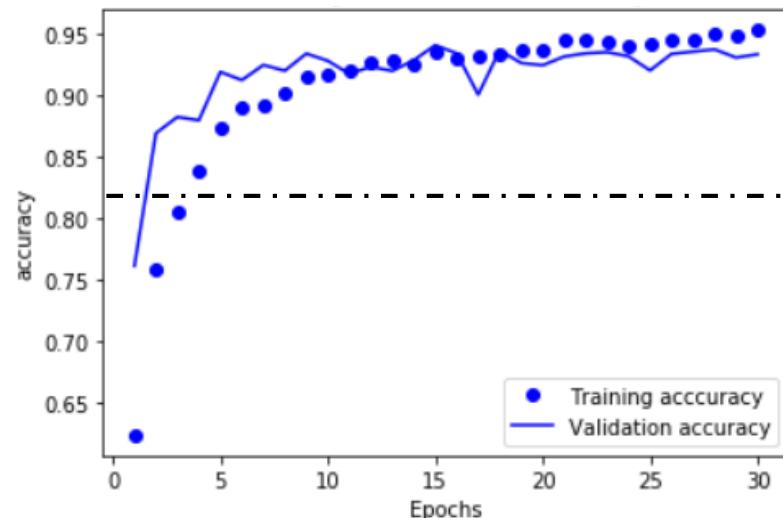
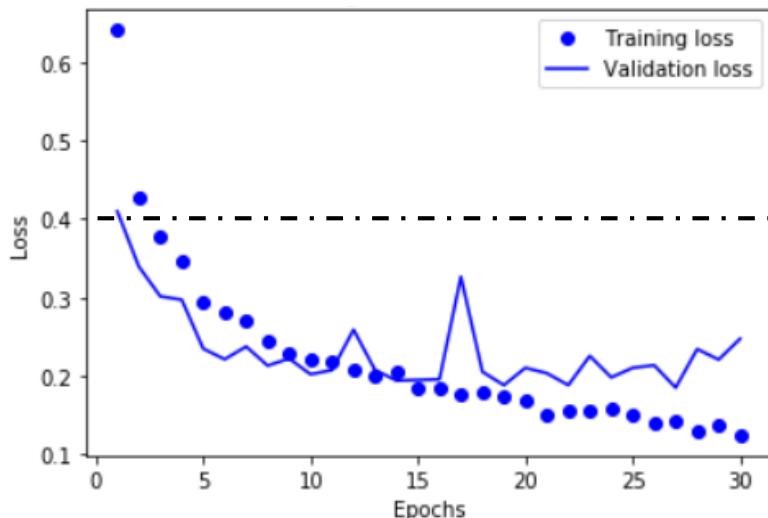
Total params: 15,928,129  
Trainable params: 1,213,441  
Non-trainable params: 14,714,688

# Fine tuning pre-trained network : efficient strategy!

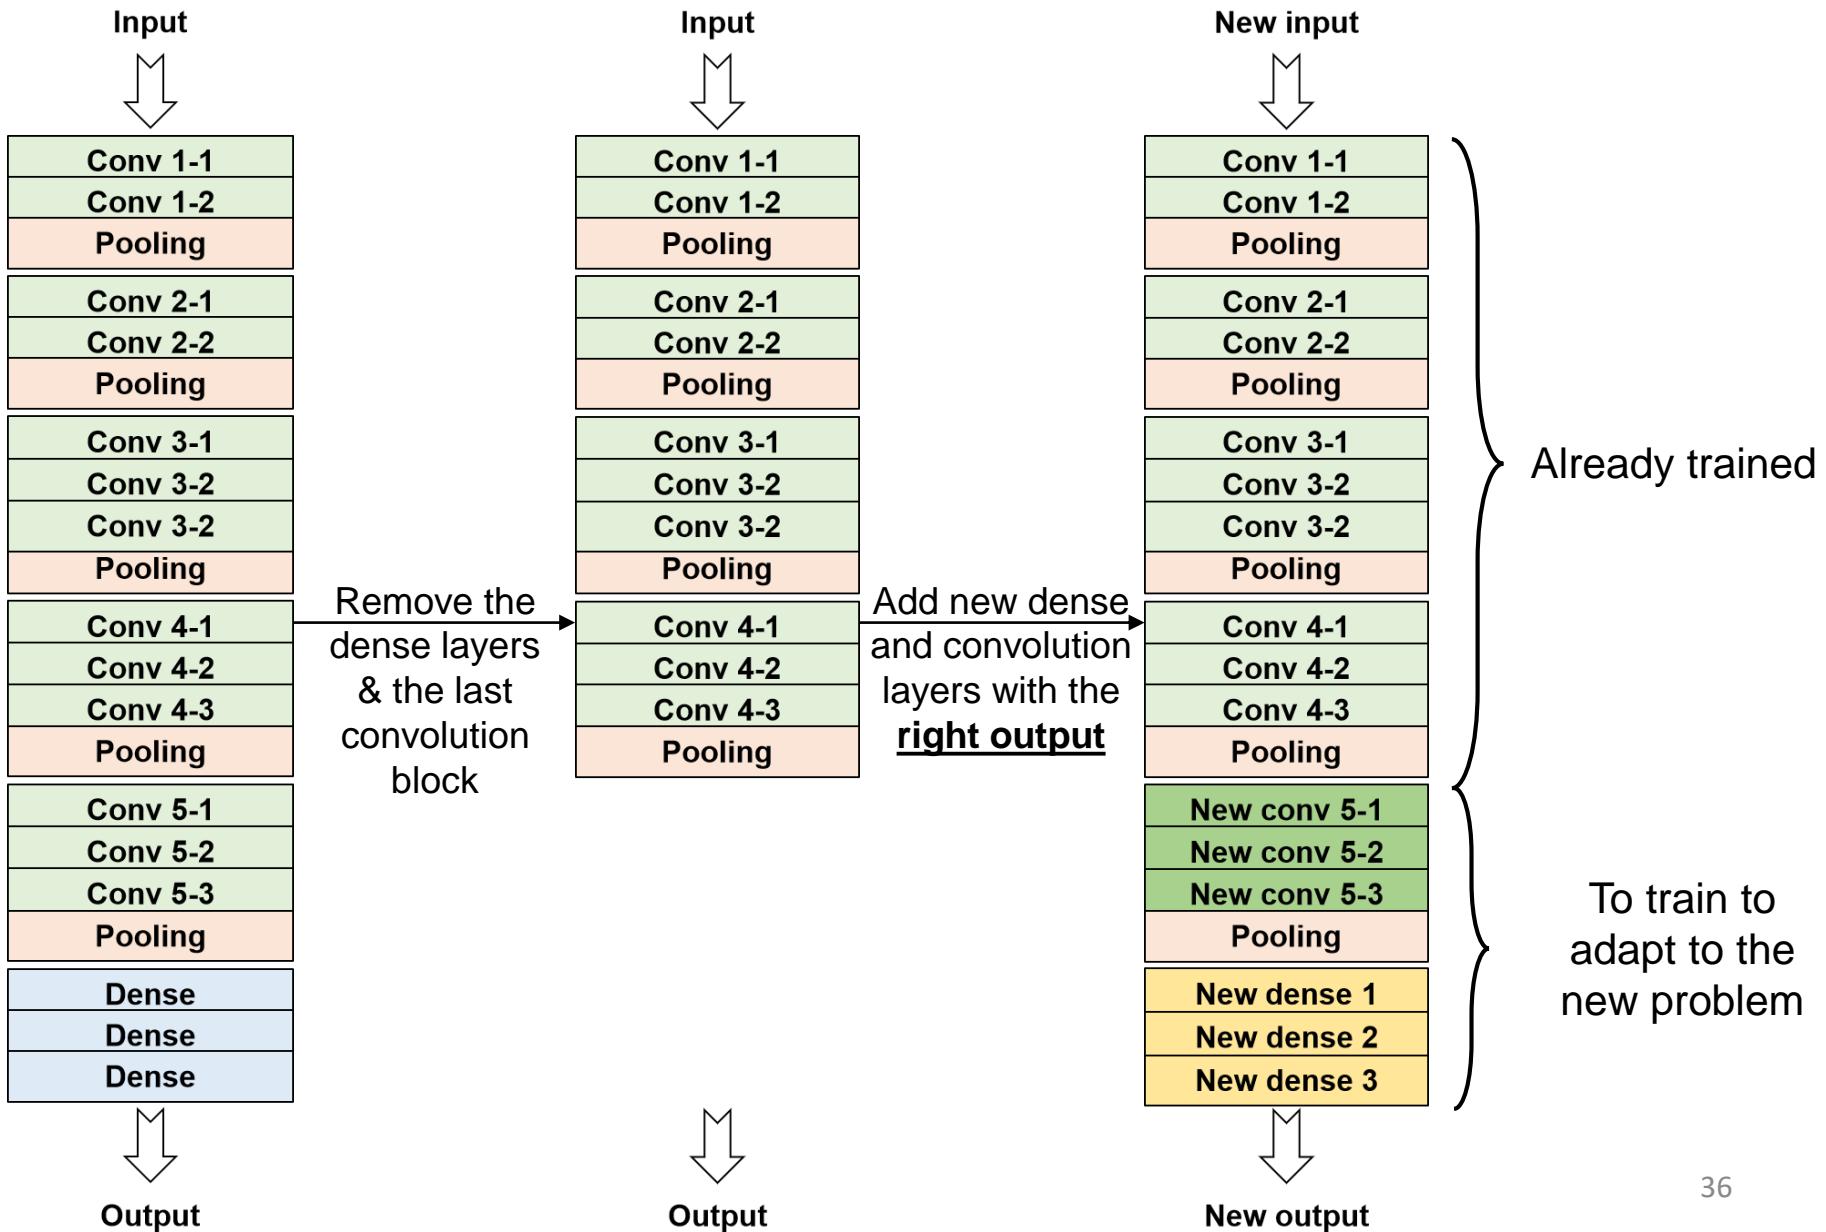
Basic convNet



Fine-tune a pre-trained network

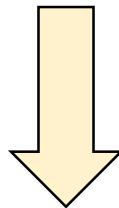


# Used a pre-trained convNet : VGG16

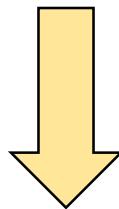


# What did we learn?

Step 1 : Clarify your problem and **assemble a dataset** for the training



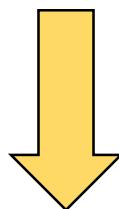
Step 2 : Try a **simple architecture** (for example MNIST) and test it



Your problem can be solved using DL?

- Training data format OK
- Measure of success
- etc.

Step 3 : Scale up the model until you **observe overfitting**



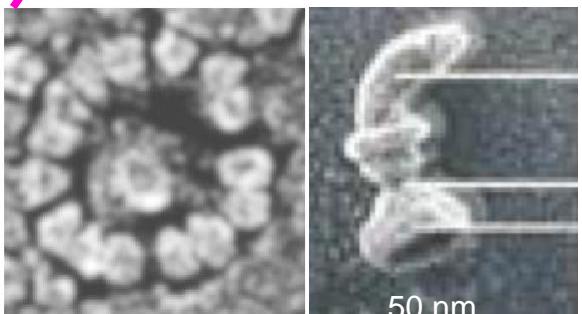
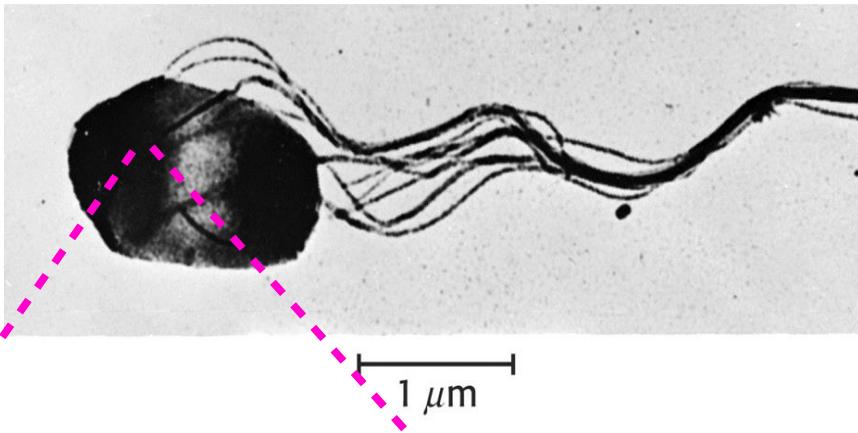
Step 4 : **Improve** your model performances

# Outline

- I. Introduction
- II. Starting with Deep Learning and Artificial Neuron Network
- III. Deep Learning for image analysis
  - i. Introduction to convolutional network
  - ii. Application : Digit recognition
  - iii. Practical case 1 : image classification
  - iv. Practical case 2 : 2D regression of bead positions**
  - v. Practical case 3 : Semantic segmentation of bacteria

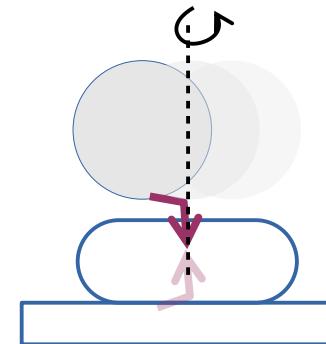
# A practical application of Deep Learning for regression

Tracking microscopic beads rotated by the flagellar motor of a bacteria

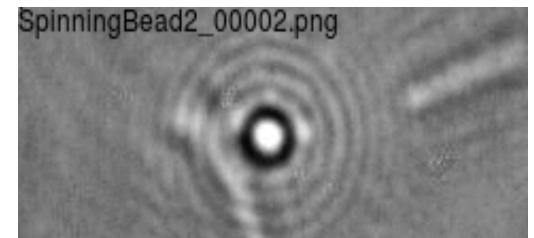


Francesco Pedaci

Ashley Nord



SpinningBead2\_00002.png



# Outline

- I. Introduction**
- II. Starting with Deep Learning and Artificial Neuron Network**
- III. Deep Learning for image analysis**
  - i. Introduction to convolutional network
  - ii. Application : Digit recognition
  - iii. Practical case 1 : image classification
  - iv. Practical case 2 : 2D regression of bead positions
  - v. Practical case 3 : Semantic segmentation of bacteria**

# Deep Learning applied to semantic segmentation



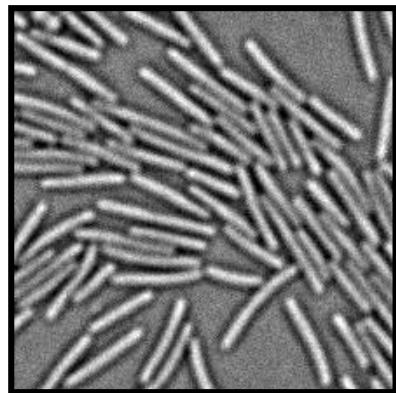
Sara Rombouts



Marcelo Nöllmann

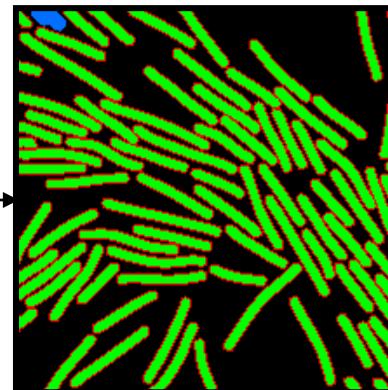


JB Fiche

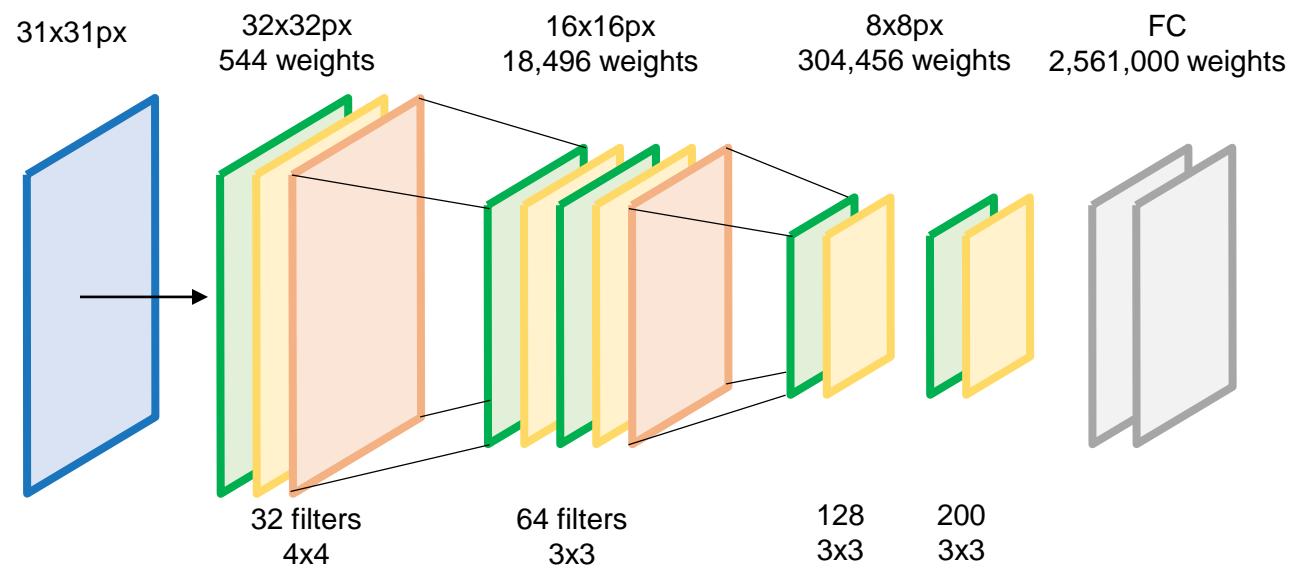
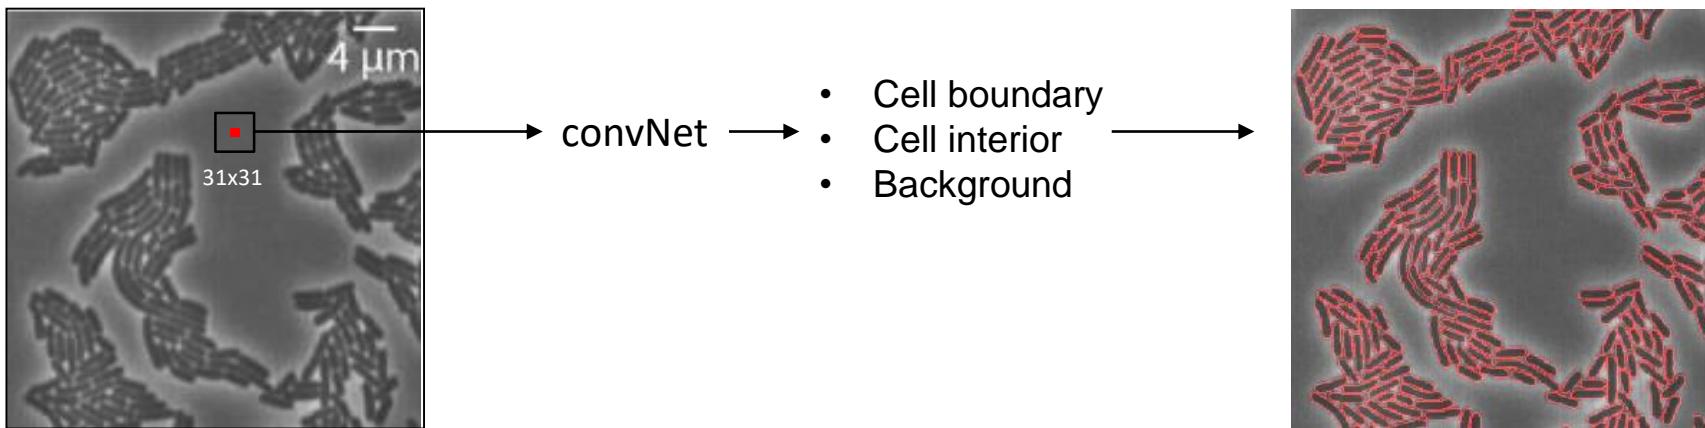


Bright field image of  
Myxococcus xanthus  
cells on an agar pad

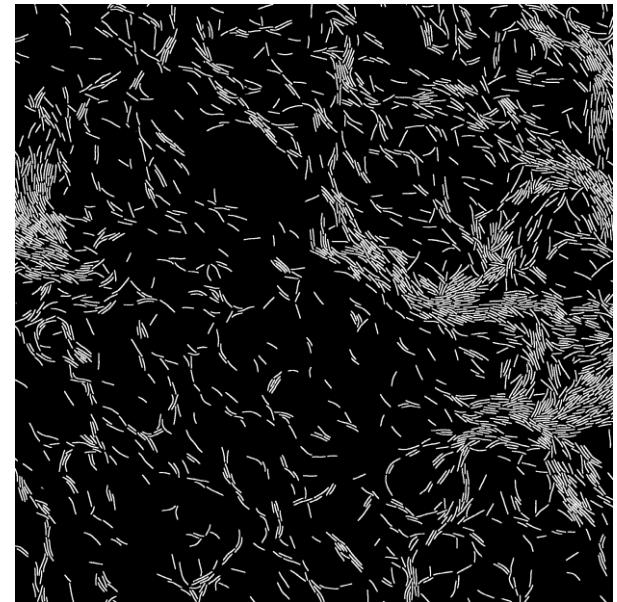
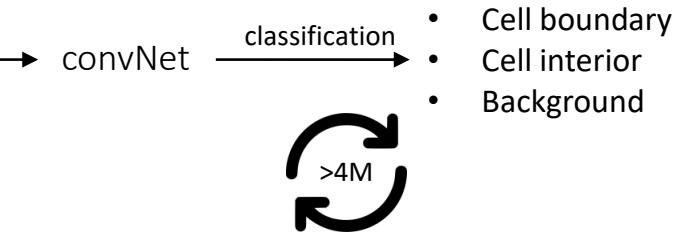
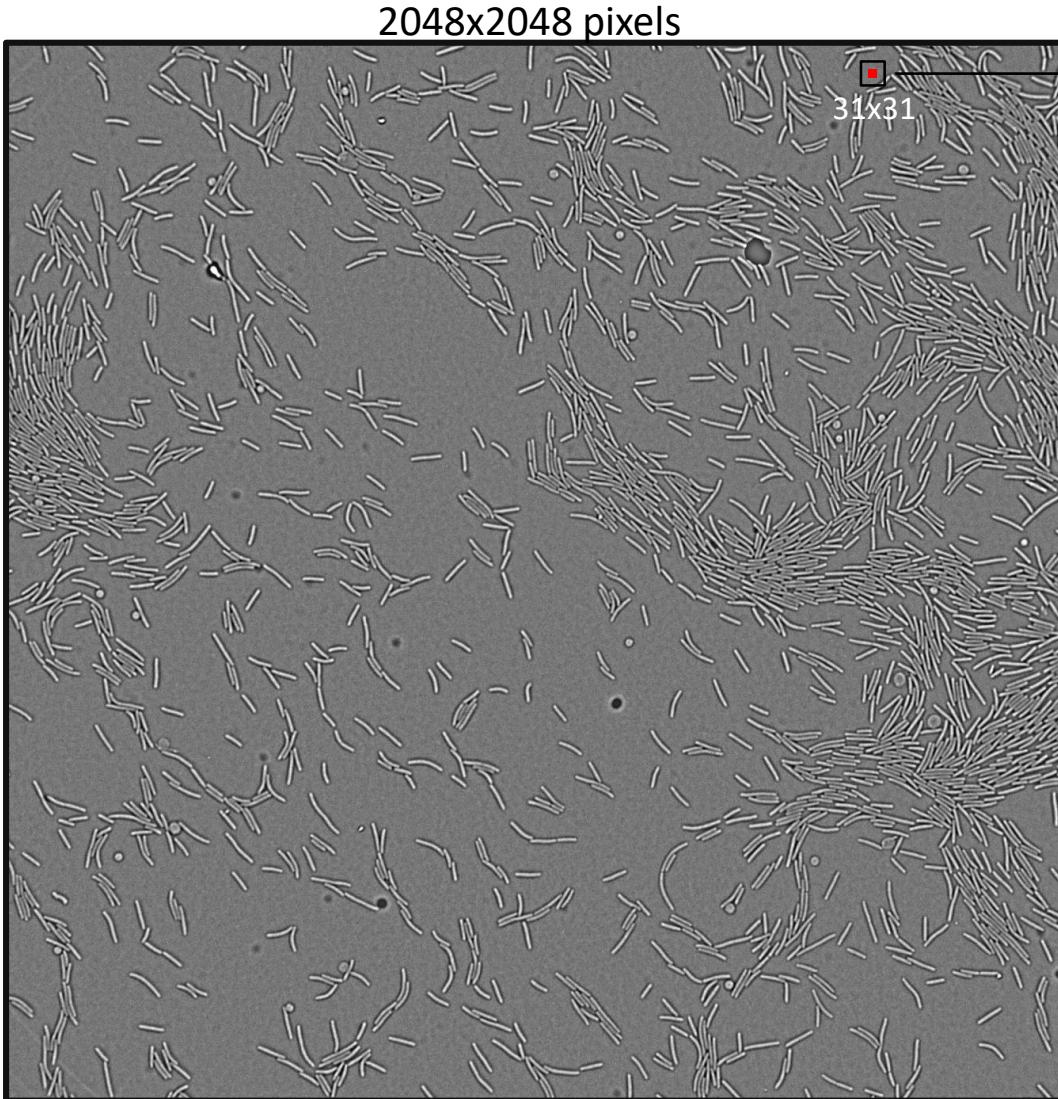
Segmentation



# ConvNet applied to cell segmentation



# Main drawback of convNet : the speed

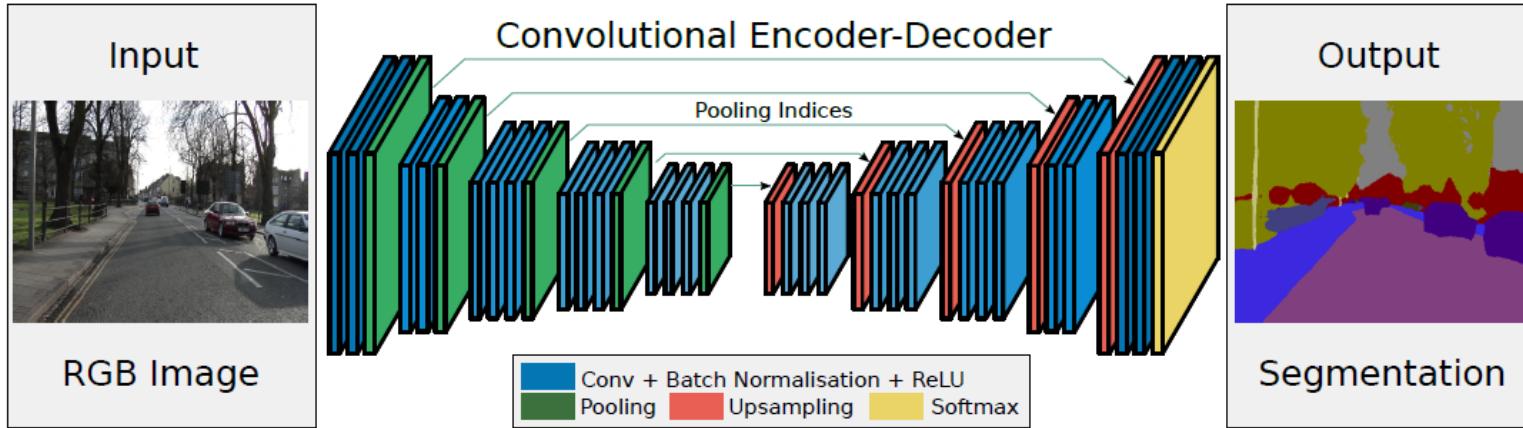


~15 minutes to segment a single image



~5 days to analyze an experiment...

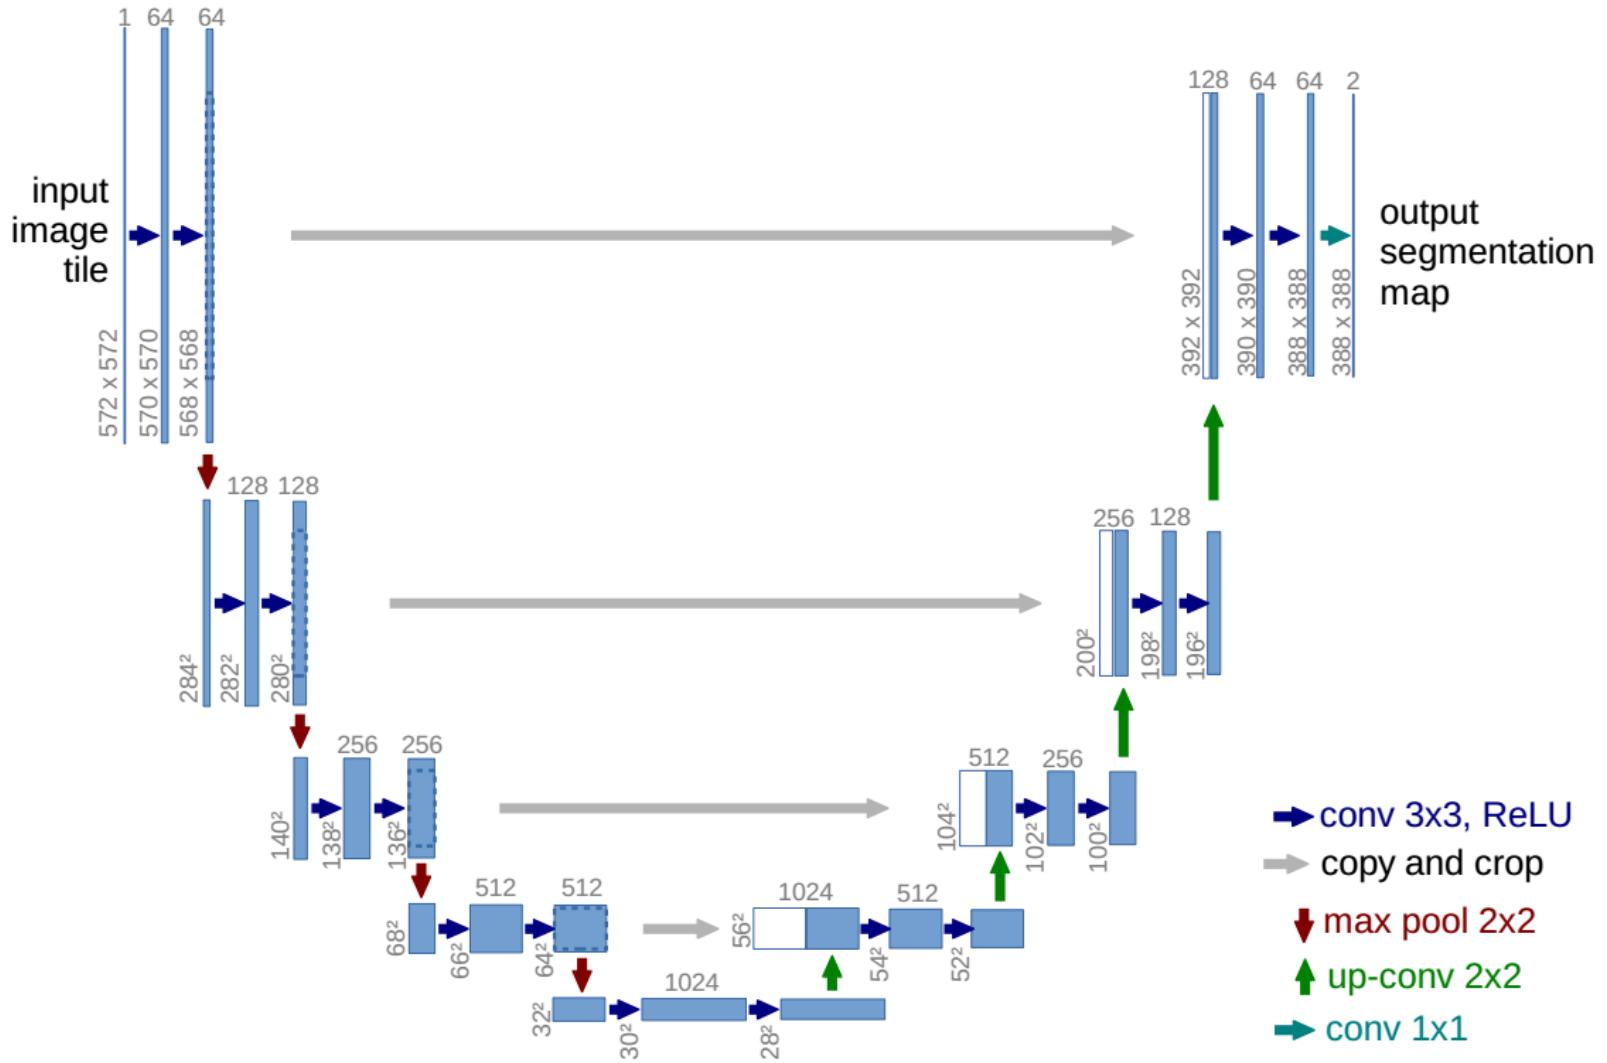
# Fully convolutional neural network : FCNN - SegNet



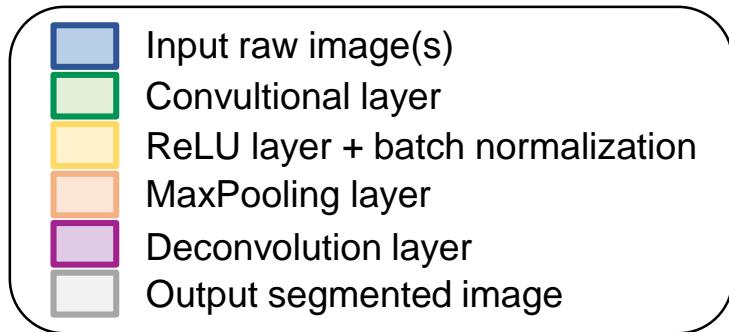
This type of neural network architecture has many advantages:

- Simple to implement
- **Faster**
- **Less heavy** on memory since no fully connected layers
- **No restriction** on the size of the input images
- The **output image is already segmented**

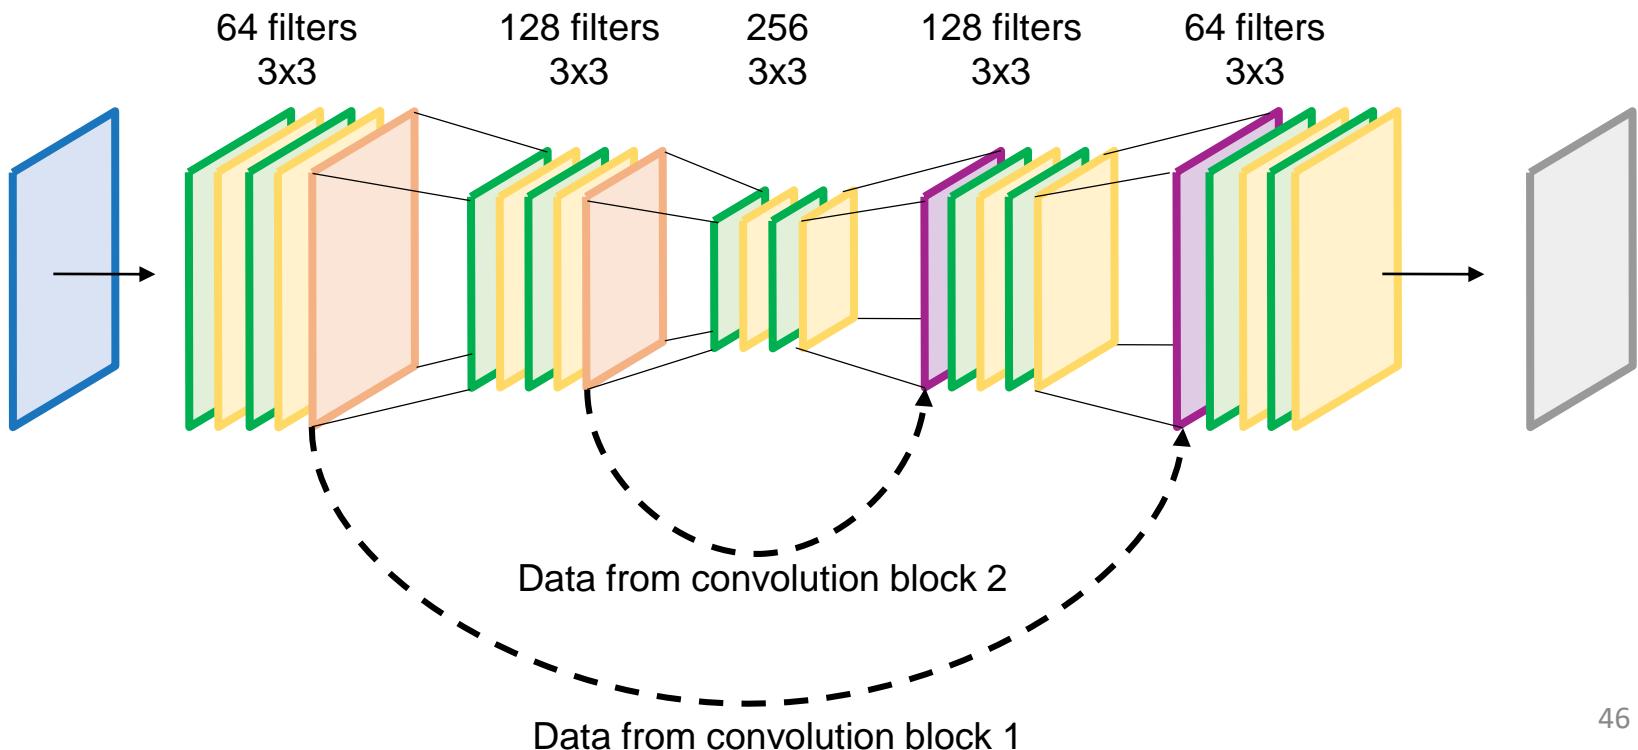
# Fully convolutional neural network : FCNN - UNet



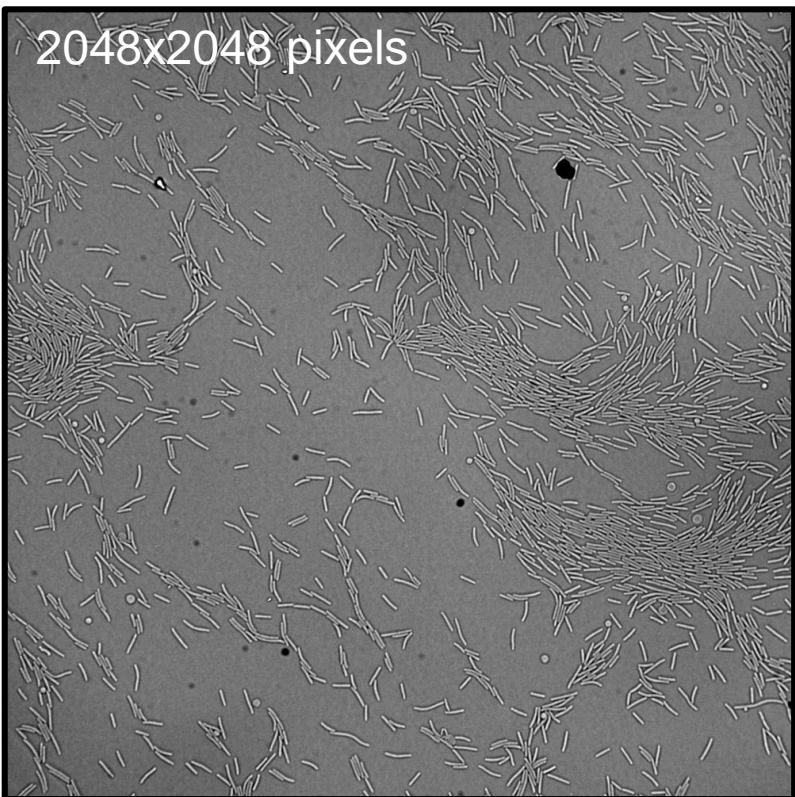
# Petit UNet ...



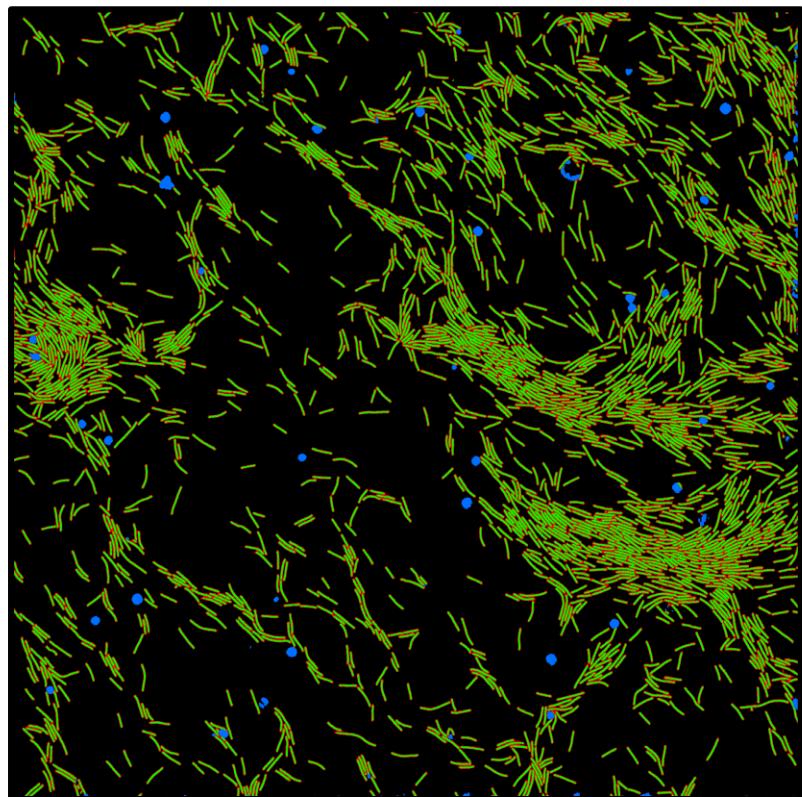
Total number of learnable : 1,864,969



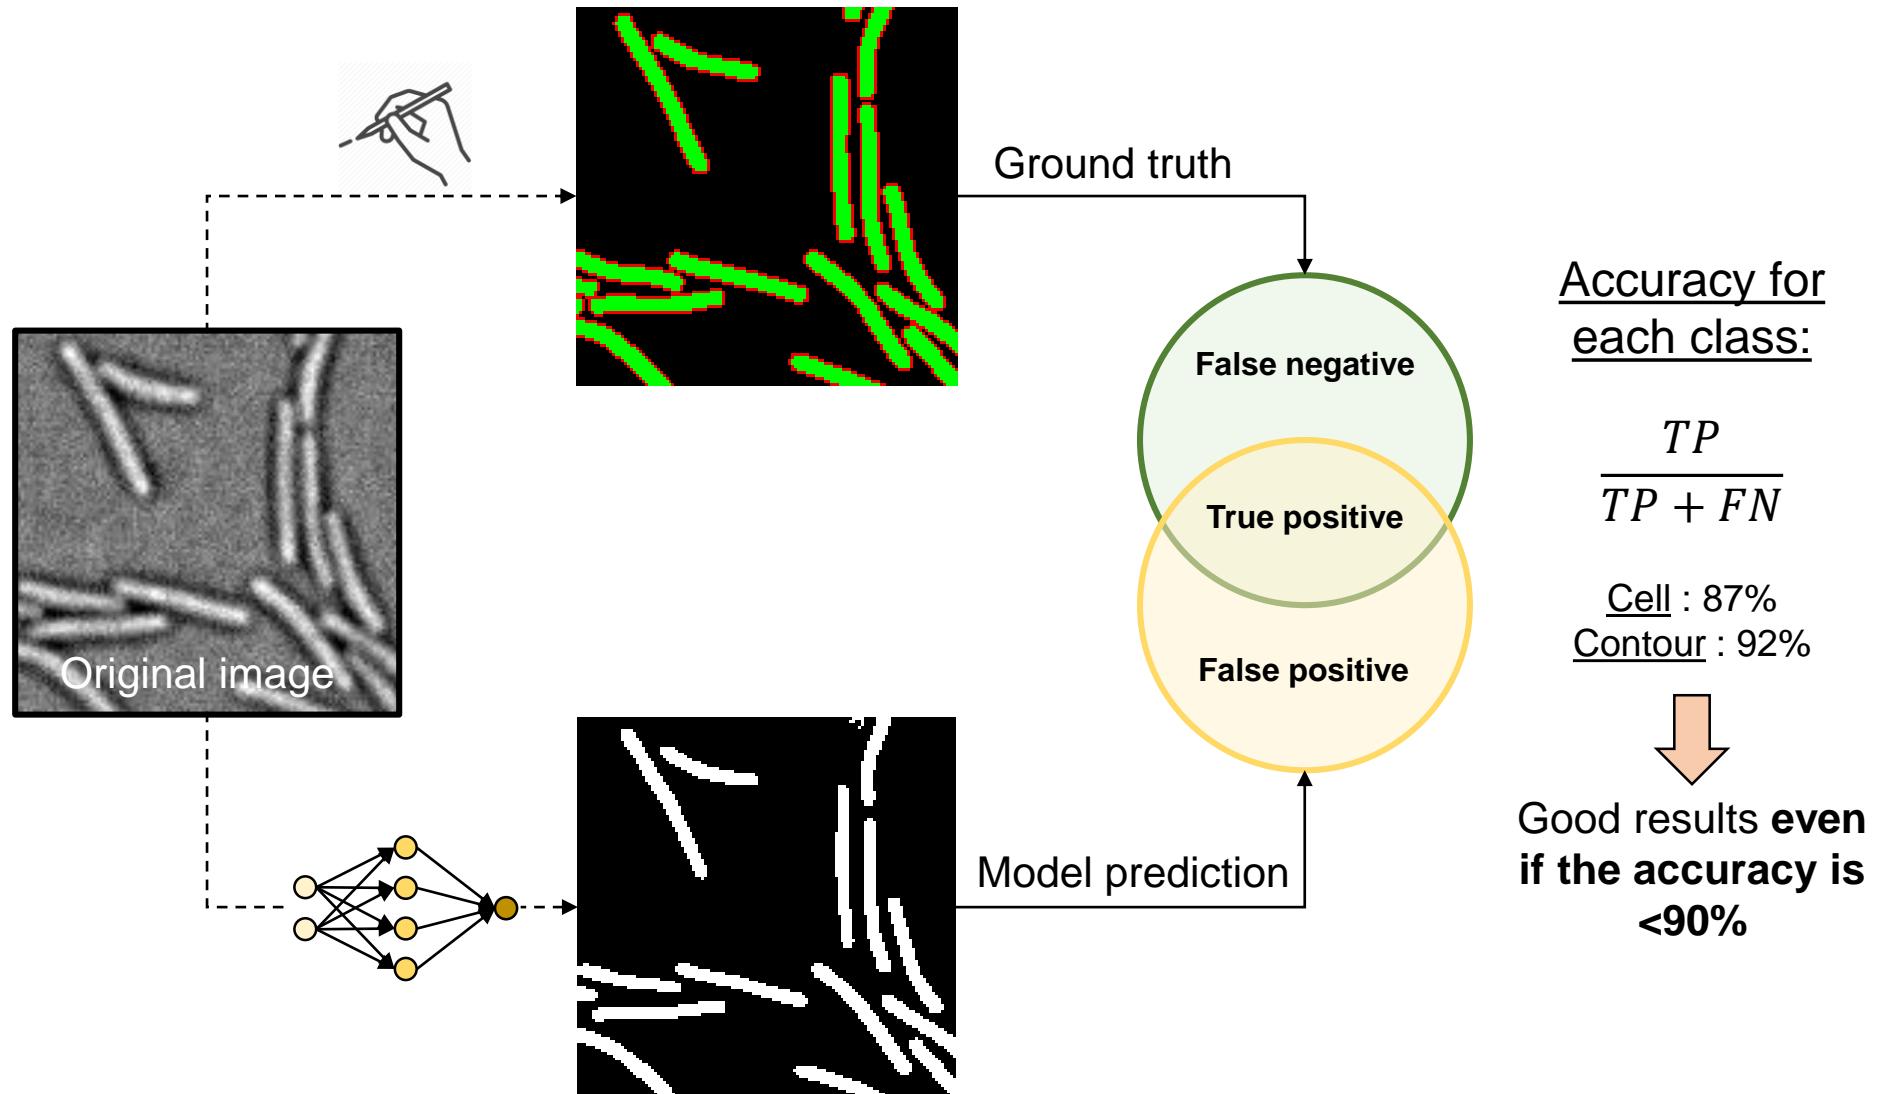
... mais costaud !



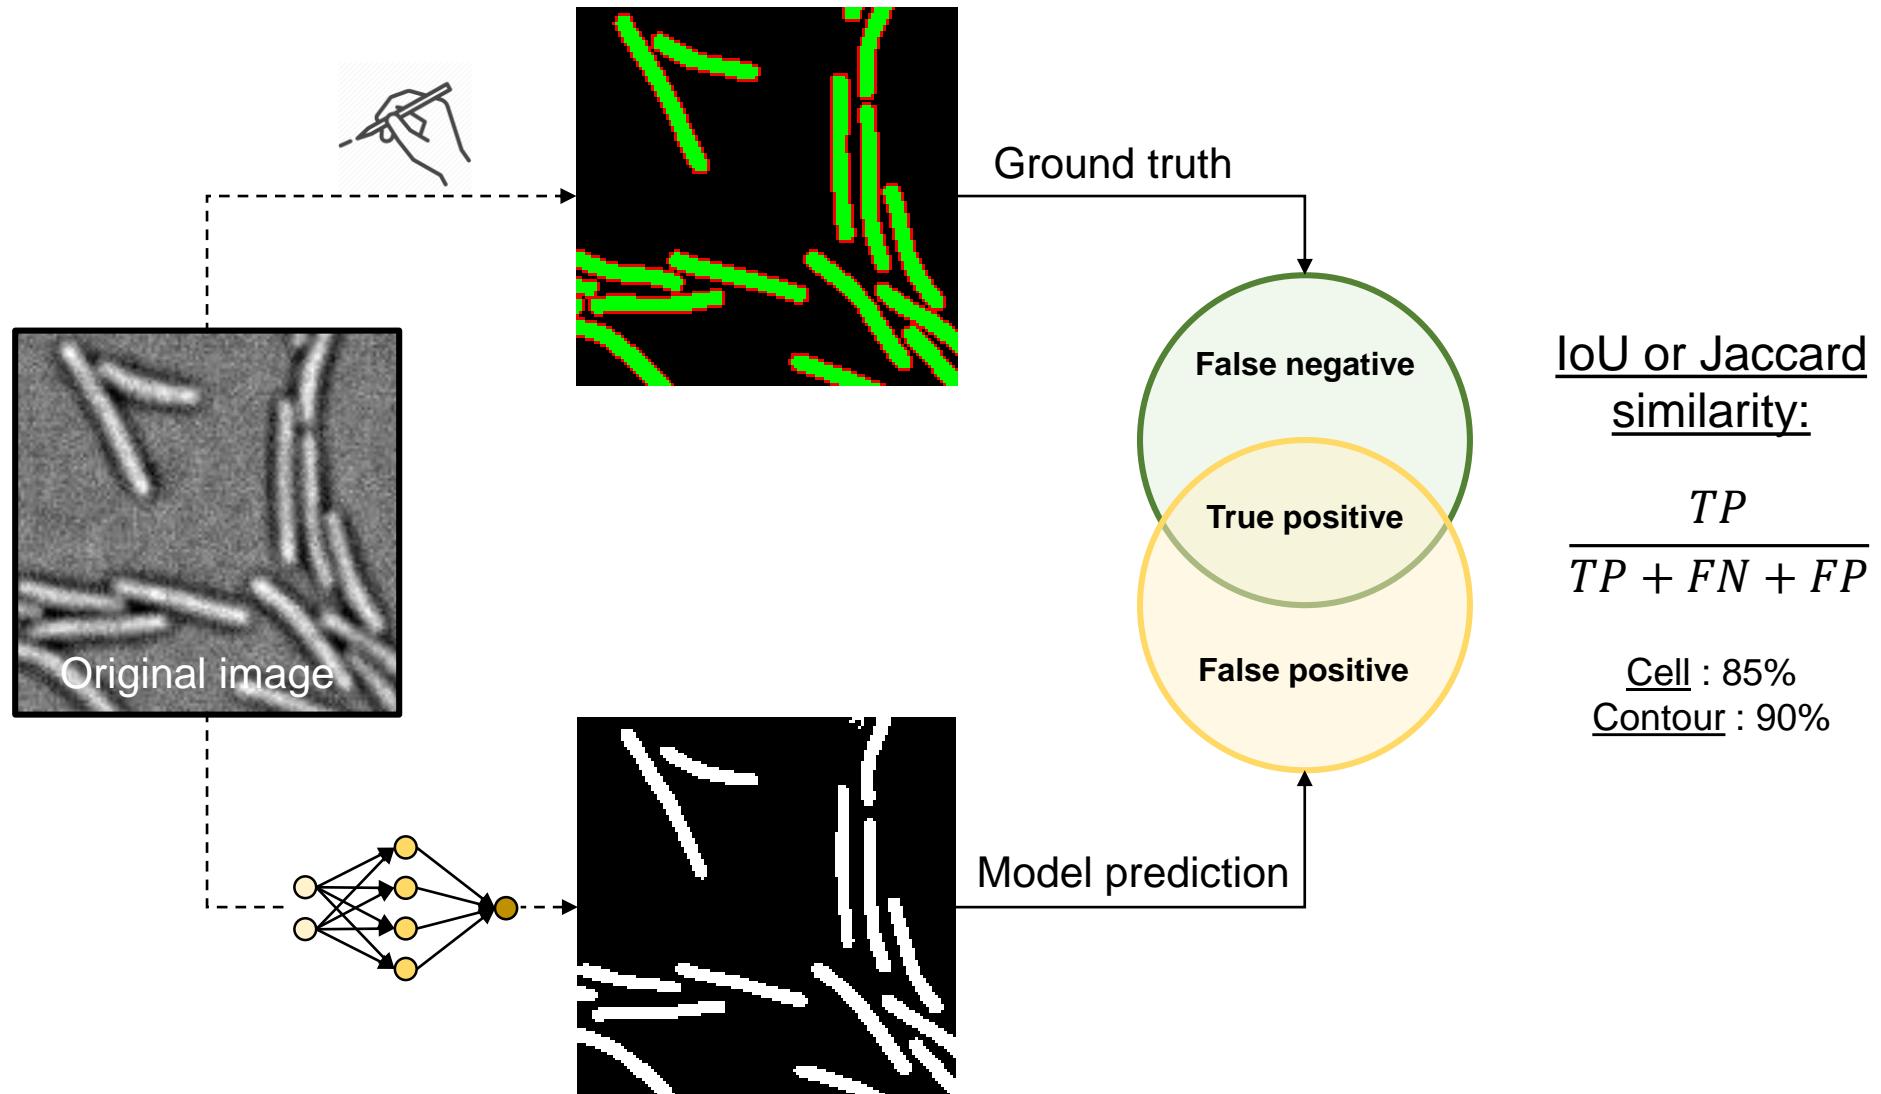
UNet  
~4s



# How to evaluate the quality of the prediction?



# How to evaluate the quality of the prediction?

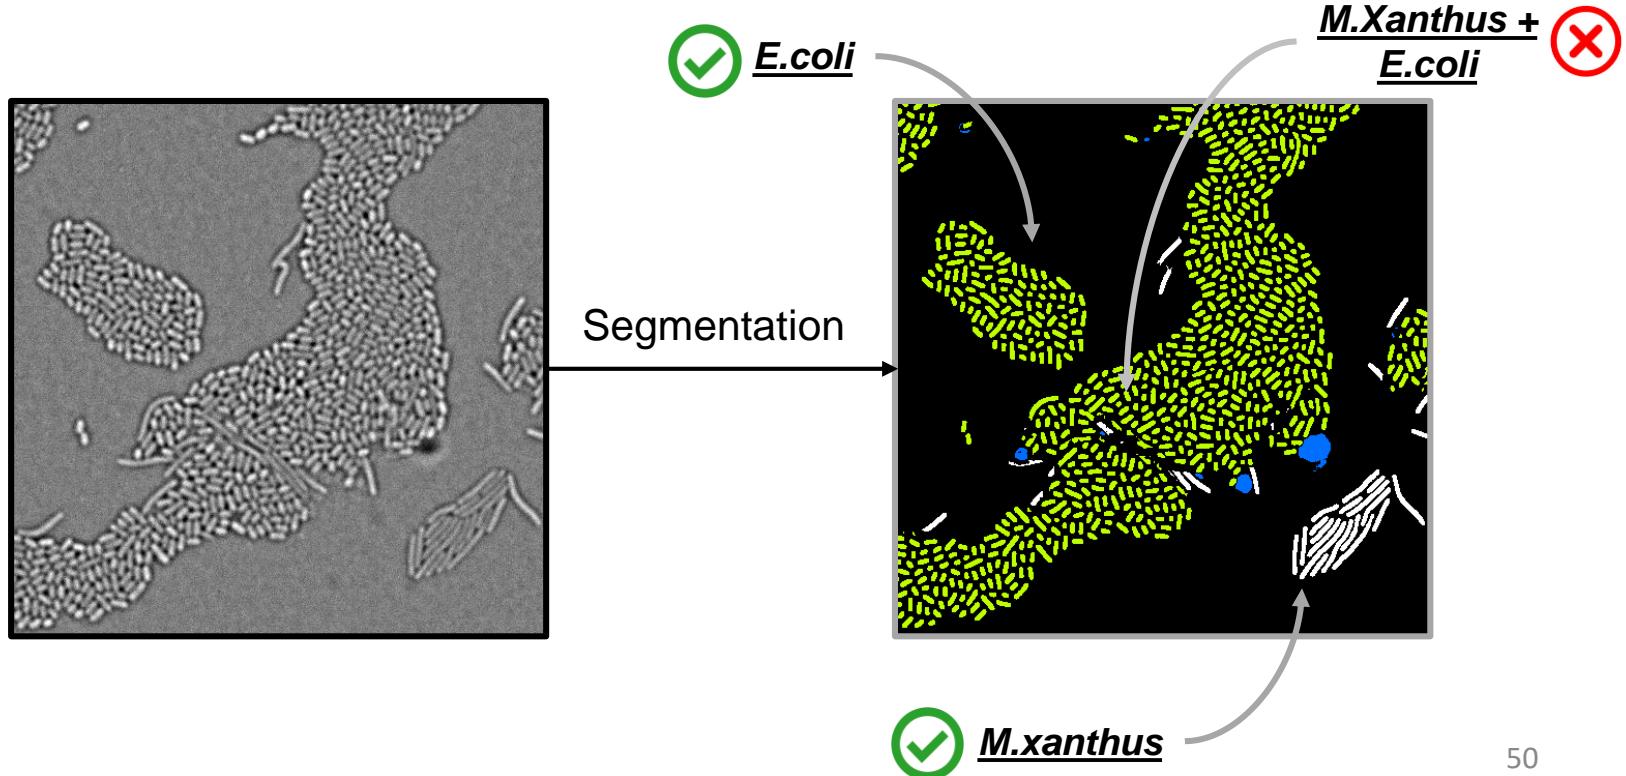


# One last VERY important point

1. Your images should be normalized:

- Average intensity equals to 0
- Standard deviation equals to 1

2. Your set of images should be as representative and diverse as possible

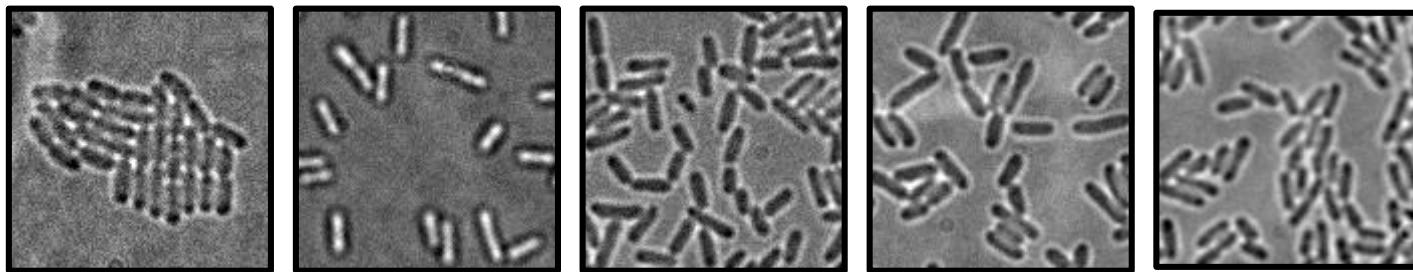


# One last VERY important point

1. Your images should be normalized:

- Average intensity equals to 0
- Standard deviation equals to 1

2. Your set of images should be **as representative and diverse as possible**



**Variation in :**

- Density of cells
- Different focal planes
- Light intensity
- etc.



**Good generalization!**

# Outline

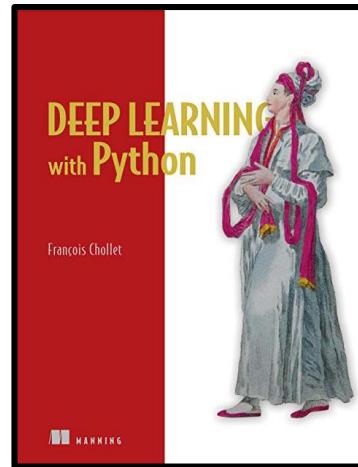
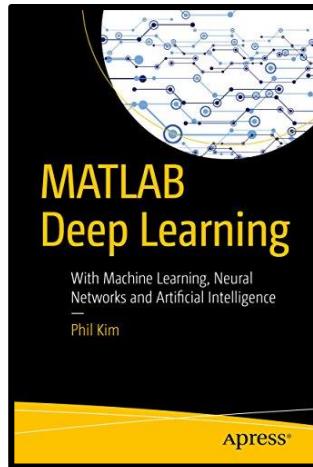
- I. Introduction
- II. Starting with Deep Learning and Artificial Neuron Network
- III. Deep Learning for image analysis
  
- IV. Conclusion...



Activation Function	Softmax
generalization	TensorFlow
overfitting	Vanishing Gradient Problem
gradient descent	VGG16
ground truth	data augmentation
hyperparameter	transfer learning
matplotlib	epoch
neuron	weights
one-hot encoding	bias
learning rate	hidden layer
Backpropagation	batch size
Categorical Cross-Entropy	classification
CNN, ConvNet	regression
Dropout	convolution
Keras	loss function
Max-Pooling	testing set
MNIST	validation set
Momentum	training set
Nonlinearity	deep
ReLU	optimizer
SGD	fully connected layer

# Bibliography & useful links

- F. Chollet – 2018 – Deep Learning with Python
- P. Kim – 2017 – Matlab Deep Learning
- I. Goodfellow, Y. Bengio, A. Courville – 2016 – Deep Learning



- Stanford lectures on Deep Learning and Machine Learning :  
<https://cs230.stanford.edu/lecture/>
- Inaugural lecture of Y. Lecun at collège de France : <https://www.college-de-france.fr/site/yann-lecun/inaugural-lecture-2016-02-04-18h00.htm>
- Awesome website for math : <https://www.3blue1brown.com/>